



Tyler Sondag  
Iowa State U.



Hridesh Rajan  
Iowa State U.

# Phase-guided Thread-to-core Assignment for Improved Utilization of Performance- Asymmetric Multi-Core Processors

International Workshop on Multicore Software Engineering


Supported in part by the US National Science Foundation under grants 06-27354 and 08-08913.

# Overview

Performance asymmetric multicores are seen as a more efficient alternative to homogeneous multicores.

**Broad Problem:** Efficient utilization of asymmetric cores

**Technical Challenge:** Match resource requirements



Asymmetric Multicore

Workload

Different shading represents varying resource requirements.

- ▶ Resource needs of threads vary at runtime.
- ▶ Target architecture may not be known statically.

**Key Insight:** Use phase behavior to reduce runtime overhead.

# Performance Asymmetric Multicores

- ▶ **What:** Cores have different characteristics (clock speed, cache size, etc.)
- ▶ **Why<sup>1</sup>:**
  - ▶ space
  - ▶ heat
  - ▶ power
  - ▶ performance-power ratio
  - ▶ parallelism




Asymmetric  
Multicore

<sup>1</sup>R. Kumar et al. ISCA '04

# Phase Behavior

- ▶ **Behavior:** resource requirements (IPC, cache, etc.)
- ▶ **Similar Behavior:** segments with similar resource usage
- ▶ **Phase:** segments of execution that exhibit similar behavior<sup>2</sup>



Phase behavior for gcc (taken from [2])


<sup>2</sup>T. Sherwood et al. ASPLOS '02

# Intuition Behind Our Solution



- ▶ **Problem:** Assign code to cores such that behavior of code matches resources of cores
- ▶ **Idea:**
  - 1 Determine sections of code that will behave in a similar way
  - 2 Knowledge of one section gives us information about all similar sections

# Approach Overview



- ▶ Idea: Apply the same thread-to-core mapping to all approximately similar sections of code
  - ❶ Statically break the program into sections of code
  - ❷ Statically determine approximate similarity between these sections
  - ❸ Dynamically monitor a section then make mapping decisions for similar section



# Program





## Ignore “small” sections









# Determine approximate similarity





# Reduce number of transition points





# Insert phase marks





# Monitor





# Run





# Run



# Monitor





# Run









# Run



# Switch to matched core




# Run on matched core



# Experimental Setup

- ▶ Hardware setup: Quad Core - 2x2.4GHz, 2x1.6GHz
- ▶ Workloads
  - ▶ 36-84 SPEC CPU2000 benchmarks
  - ▶ constant workload size
- ▶ Compare to standard Linux assignment

# Overall



**Best Result:** Interval technique, min. size 45 instructions

# Previous Work

Falls into two categories

- ▶ Asymmetry-aware scheduler<sup>3</sup>
  - ▶ high monitoring overhead
  - ▶ requires OS modification
- ▶ Improved load balancing<sup>45</sup>
  - ▶ ignores behavior - may cause inefficient utilization
  - ▶ requires OS modification

---

<sup>3</sup>R. Kumar et al. ISCA '04

<sup>4</sup>T. Li et al. SC '07

<sup>5</sup>M. Becchi et al. CF '06

# Conclusion

- ▶ Performance asymmetric multicores are a beneficial class of processors.
- ▶ Problem: Techniques to effectively assign threads to cores are still needed.
- ▶ **Solution: Use phase behavior to reduce dynamic overhead.**
  - ▶ Programmer oblivious
  - ▶ Automatic
  - ▶ Negligible overhead
  - ▶ Transparent deployment

# Questions

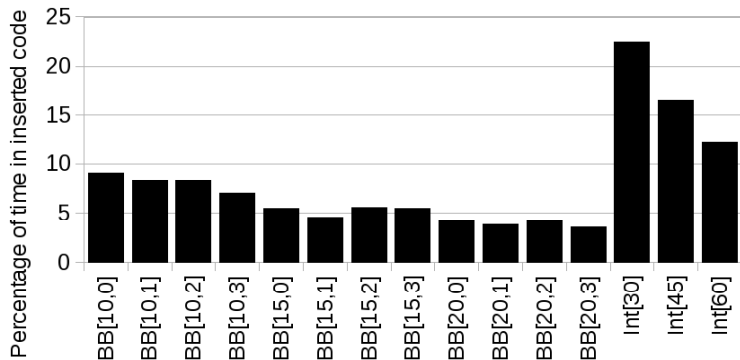
# Questions?



# Experimental Setup


- ▶ Hardware setup: Quad Core - 2x2.4GHz, 2x1.6GHz
- ▶ Software setup
  - ▶ Static analysis/instrumentation: our framework based on GNU Binutils
  - ▶ Runtime Performance monitoring: PAPI, perfmon2
  - ▶ Core switching: affinity calls built-in to kernel
  - ▶ Workloads
    - ▶ 36-84 SPEC CPU2000 benchmarks
    - ▶ constant workload size
- ▶ Compare to standard Linux assignment

# Overheads (Time)




**BB[x, y]:** Basic block technique, min. block size: x,  
Look-ahead: y. **Int[x]:** interval technique, min. interval size: x

# Throughput Improvement (Instructions Executed)




Left: Interval technique, Right: Basic block technique


# Speedup vs Fairness



# Speedup vs Overhead



# Speedup vs Throughput



# Determining program behavior

Falls into two categories

- ▶ Techniques using execution traces
- ▶ Purely dynamic techniques

# Execution Traces

- ▶ **Benefits:**
  - ▶ Very accurate since actual performance is known
  - ▶ Low dynamic overhead since no monitoring is required
- ▶ **Limitations:**
  - ▶ Requires sample input set to be developed
  - ▶ Run entire program to create execution trace
  - ▶ What about sections of code not covered by sample input?
  - ▶ Do different inputs result in different behavior?



# Purely Dynamic

- ▶ Benefits:

- ▶ Does not require sample input sets
- ▶ No need for execution trace
- ▶ Does not monitor the whole program

- ▶ Limitations:

- ▶ Decisions for future code are made based on past code
- ▶ Higher dynamic overhead since we must monitor periodically throughout the entire execution

# Static Phase Marking

- ▶ Predict similarity between sections of code
- ▶ Insert phase marks on type transitions if determined beneficial
  - ▶ Basic blocks with look-ahead
  - ▶ Intervals

## Phase marks

- ▶ Dynamic analysis code
  - ▶ Monitor code if no mapping is unknown
  - ▶ Switch cores if mapping is known
- ▶ Type information

# Asymmetry Aware Scheduler

- ▶ **What:** Scheduler assigns threads to well matched cores
- ▶ **Benefits:**
  - ▶ Very accurate since based on actual performance
  - ▶ Makes system wide decisions
  - ▶ Programs switch cores as behavior changes
- ▶ **Limitations:**
  - ▶ Monitoring is required throughout entire execution
  - ▶ Decisions for future execution are based on past behavior
  - ▶ Requires OS modification

# Improved Load Balancing

- ▶ **What:** “Fast” cores get more processes or round-robin
- ▶ **Benefits:**
  - ▶ Low overhead: does not monitor execution
  - ▶ System wide decision making
- ▶ **Limitations:**
  - ▶ Aimed at fairness
  - ▶ Ignores behavior: “Fast” programs may be on “slow” cores
  - ▶ Requires OS modification

# Intuition Behind Our Solution

- ▶ **Problem:** Assign code sections to cores such that behavior of code matches resources of cores
- ▶ **Idea:** If we can determine that sections of code will behave in a similar way, knowledge of one section gives us information about all similar sections.
- ▶ **Advantages** of this approach
  - ▶ Only need to monitor a small amount of code dynamically
  - ▶ No need to predict the actual behavior, just similarity.
  - ▶ Considers changes in program behavior
  - ▶ No knowledge of target machine required
  - ▶ No need for execution traces or sample inputs
  - ▶ Automatic - transparent to programmer and end user