

# Bytecodes Meet Combinators

John Rose

Architect, Da Vinci Machine Project

Sun Microsystems

VMIL 2009: 3rd Workshop on Virtual Machines and Intermediate Languages  
(OOPSLA, Orlando, Florida October 25, 2009)

<http://cr.openjdk.java.net/~jrose/pres/200910-BcsMeetCmbs.pdf>

# Overview

- The JVM is useful for more languages than Java
- Bytecodes have been successful so far, but are showing some “pain points”
- Composable “function pointers” fill in a key gap

# JVM Architecture

agents & plugins
start/exit/abort hooks
debugging
profiling
instrumentation & monitoring
serviceability

memory systems
compiled code cache
managed object heap
GCs: parallel scavenge, concurrent mark/sweep, regionalized, ...

user code & library code
classfiles
JNI methods
resource files

misc. JVM primitive APIs
JVM configuration
security, access managers
I/O & OS interfaces
Java reflection APIs
unsafe operations

class loading
bytecode verifier
class linker
bytecode interpreter
ClassLoader API
Java native interface (JNI)

threads
stack walker
locks & safepoints
stack frame transformer

dynamic compilation
front-end/optimizer/back-end
type & frequency profile
dependency tracking
(re-/de-)compilation policy

operating system	hardware
------------------	----------

# “Java is slow because it is interpreted”

- Early implementations of the JVM executed bytecode with an interpreter **[slow]**





# “Java is fast because it runs on a VM”



- Major breakthrough was the advent of “Just In Time” compilers [fast]
  - > Compile from bytecode to machine code at runtime
  - > Optimize using information *available at runtime only*
- Simplifies static compilers
  - > javac and ecj generate “dumb” bytecode and trust the JVM to optimize
  - > Optimization is *real*...  
...even if it is invisible

# HotSpot optimizations

## compiler tactics

- delayed compilation
- tiered compilation
- on-stack replacement
- delayed reoptimization
- program dependence graph representation
- static single assignment representation

## proof-based techniques

- exact type inference
- memory value inference
- memory value tracking
- constant folding
- reassociation
- operator strength reduction
- null check elimination
- type test strength reduction
- type test elimination
- algebraic simplification
- common subexpression elimination
- integer range typing

## flow-sensitive rewrites

- conditional constant propagation
- dominating test detection
- flow-carried type narrowing
- dead code elimination

## language-specific techniques

- class hierarchy analysis
- devirtualization
- symbolic constant propagation
- autobox elimination
- escape analysis
- lock elision
- lock fusion
- de-reflection

## speculative (profile-based) techniques

- optimistic nullness assertions
- optimistic type assertions
- optimistic type strengthening
- optimistic array length strengthening
- untaken branch pruning
- optimistic N-morphic inlining
- branch frequency prediction
- call frequency prediction

## memory and placement transformation

- expression hoisting
- expression sinking
- redundant store elimination
- adjacent store fusion
- card-mark elimination
- merge-point splitting

## loop transformations

- loop unrolling
- loop peeling
- safepoint elimination
- iteration range splitting
- range check elimination
- loop vectorization

## global code shaping

- inlining (graph integration)
- global code motion
- heat-based code layout
- switch balancing
- throw inlining

## control flow graph transformation

- local code scheduling
- local code bundling
- delay slot filling
- graph-coloring register allocation
- linear scan register allocation
- live range splitting
- copy coalescing
- constant splitting
- copy removal
- address mode matching
- instruction peepholing
- DFA-based code generator

# Feedback multiplies optimizations

- On-line profiling and CHA produces information
- ...which lets the JIT ignore unused paths
- ...and helps the JIT sharpen types on hot paths
- ...which allows calls to be devirtualized
- ...allowing them to be inlined
- ...expanding an ever-widening optimization horizon

*Result: Large nmethods containing tightly optimized machine code for hundreds of inlined calls.*

# Why should Java have all the fun?

These optimizations apply very well elsewhere

- Static languages: JavaFX, Scala, ...
- Dynamic languages: JRuby, JavaScript, Python, ...
- Functional languages: Clojure, Scala (again)




“Ruby is slow because it is interpreted”

***NOT!***



# The Great Ruby Shootout 2008




2.00  
means  
“twice  
as fast”

0.50  
means  
“half  
the  
speed”

# Languages ♥ The JVM



# Different *kinds* of languages



Data courtesy of TIOBE: [www.tiobe.com](http://www.tiobe.com)

# JVM Specification, 1997

The Java Virtual Machine knows nothing about the Java programming language, only of a particular binary format, the class file format.

A class file contains Java Virtual Machine instructions (or bytecodes) and a symbol table, as well as other ancillary information.

Any language with functionality that can be expressed in terms of a valid class file can be hosted by the Java virtual machine.

Attracted by a generally available, machine-independent platform, implementors of other languages are turning to the Java Virtual Machine as a delivery vehicle for their languages.

**In the future, we will consider bounded extensions to the Java Virtual Machine to provide better support for other languages.**

# Predicting the future is not easy...

- Hardware: Asymmetric MC? RISC? VPU ? Fibers?
- Memory: NUMA? Flash tier? (S)TM? Functional?
- Systems: Central? Mobile? Embedded? Realtime?
- Languages: Parallel? Reactive? Dynamic?
- Programming: Script-based? Tool-based?
- Java Franchise: Java++? Other Java companies?



# Key bet: Run dynamic languages

- Scripting is here to stay.  
PHP, JavaScript, Ruby, Python  
Key value: Flexibility (informal notations)
- Late binding is getting cheaper  
(late composition getting more common)
- The JVM must interoperate with such languages.
- Tactic: JSR 292, the “invokedynamic” bytecode.

# Bytecode Strengths (historical)


- Compact
- Portable
- Verifiable
- Composable (linkage through names, everywhere)
- Flexible: *Both* interpretable and compilable

=> many machine cycles run via JVM bytecodes

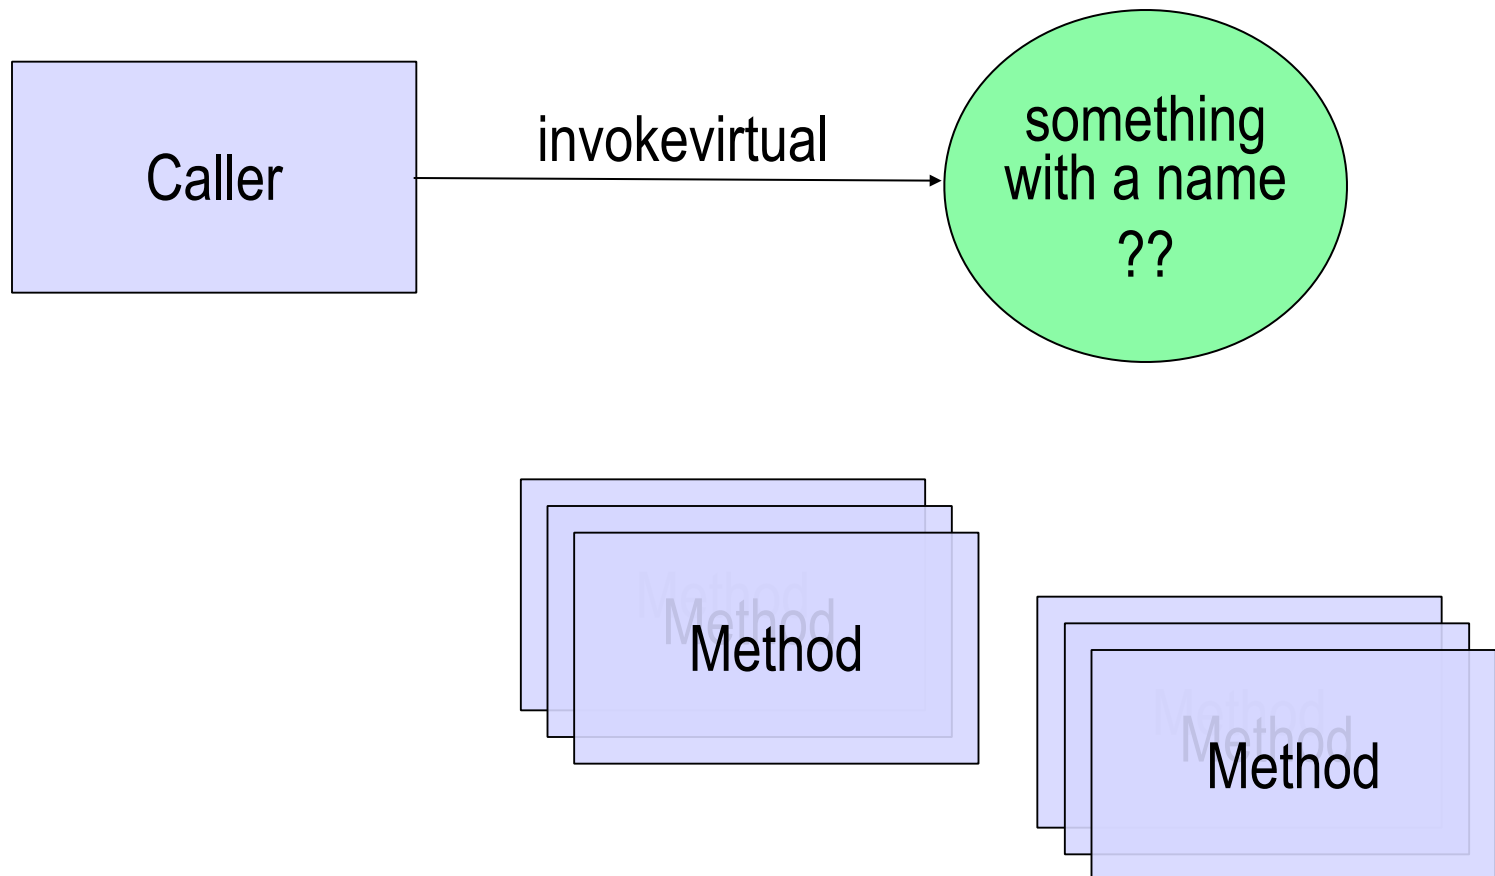
# Bytecode Pain Points


- Designed for one language (starts with a “J”)
- No “join points” other than name resolution
- Name resolution is not (very) programmable
- Changes require reassembly
- => reverification, recompilation, reoptimization
- Small devices object to bytecode weaving

# Key composition: Invoke a method



# Initial state is unlinked







## Digression: Details of Method Calls

The next group of slides, headlined “JavaOne”, are taken from Brian Goetz’s and John’s JavaOne 2009 talk, [Towards a Renaissance VM](#). These slides were omitted from the actual VMIL presentation.

The full talk is available:

<http://cr.openjdk.java.net/~jrose/pres/200906-RenaisVM.pdf>

The related talk “JSR 292 Cookbook” is here:

<http://cr.openjdk.java.net/~jrose/pres/200906-Cookbook.htm>

## The deal with method calls (in one slide)

- > Calling a method is cheap (VMs can even inline!)
- > Selecting the right target method can be costly
  - > Static languages do most of their method selection at compile time (e.g., `System.out.println(x)`)
    - > Single-dispatch on receiver type is left for runtime
  - > Dynamic langs do almost none at compile-time
    - > But it would be nice to not have to re-do method selection for *every single invocation*
- > Each language has its own ideas about linkage
  - > The VM enforces static rules of naming and linkage
    - > Language runtimes want to decide (& re-decide) linkage

# What's in a method call?

- > Naming — using a symbolic name
- > Linking — reaching out somewhere else
- > Selecting — deciding which one to call
- > Adapting — agreeing on calling conventions
- > *(...and finally, a parameterized control transfer)*

## A connection from caller A to target B

- > Including naming, linking, selecting, adapting:
- > ...where B might be known to A only by a name
- > ...and A and B might be far apart
- > ...and B might depend on arguments passed by A
- > ...and a correct call to B might require adaptations
- > *(After everything is decided, A jumps to B's code.)*

## Example: Fully static invocation

- > For this source code

```
String s = System.getProperty("java.home");
```

The compiled byte code looks like

```
0:   ldc #2           //String "java.home"
2:   invokestatic #3   //Method java/lang/System.getProperty:
                          (Ljava/lang/String;)Ljava/lang/String;
5:   astore_1
```

## Example: Fully static invocation

- > For this source code

```
String s = System.getProperty("java.home");
```

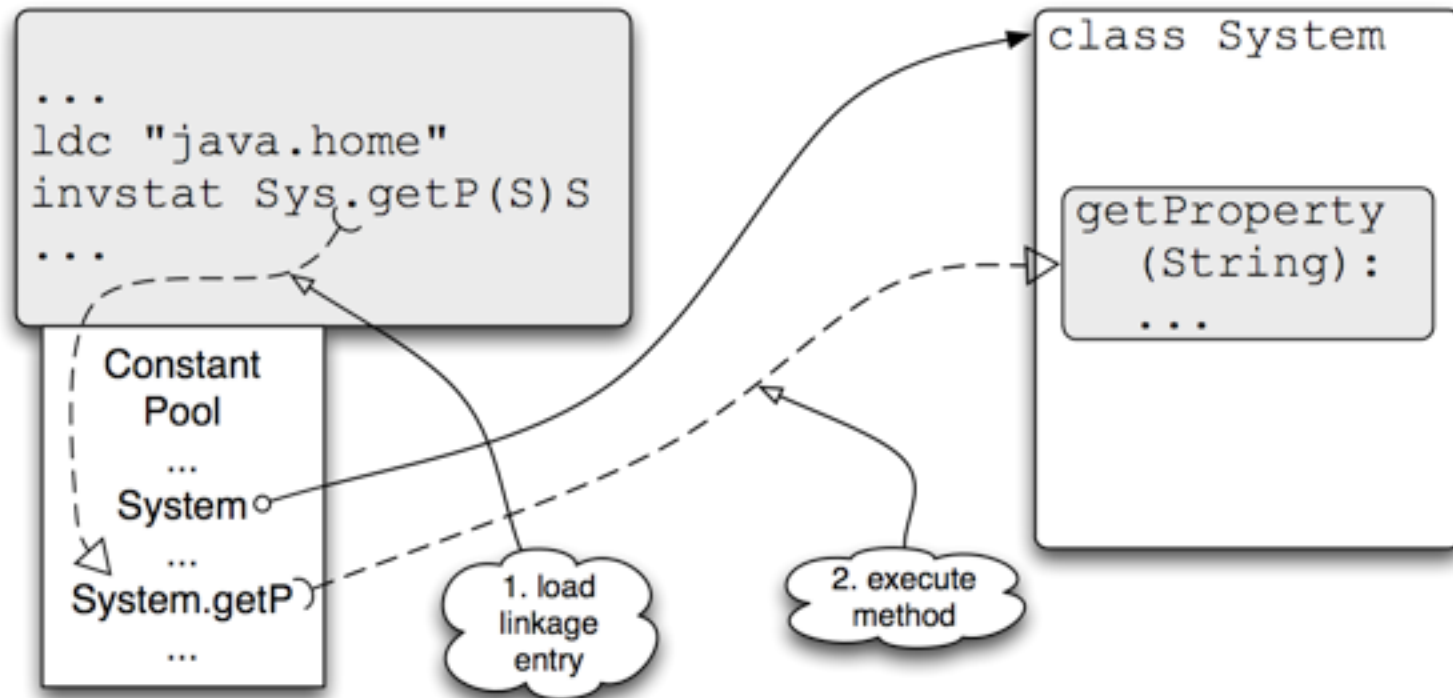
The compiled byte code looks like

```
0:   ldc #2           //String "java.home"
2:   invokestatic #3   //Method java/lang/System.getProperty:
                        (Ljava/lang/String;)Ljava/lang/String;
5:   astore_1
```

- | Names are embedded in the bytecode
- | Linking handled by the JVM with fixed Java rules
- | Target method selection is not dynamic at all
- | No adaptation: Signatures must match exactly



## How the VM sees it:



&gt;

(Note: This implementation is typical; VMs vary.)

# Example: Class-based single dispatch

- > For this source code

```
//PrintStream out = System.out;  
out.println("Hello World");
```

- > The compiled byte code looks like

```
4:    aload 1  
5:    ldc #2          //String "Hello World"  
7:    invokevirtual #4 //Method java/io/PrintStream.println:  
                                (Ljava/lang/String;)V
```

## Example: Class-based single dispatch

- > For this source code

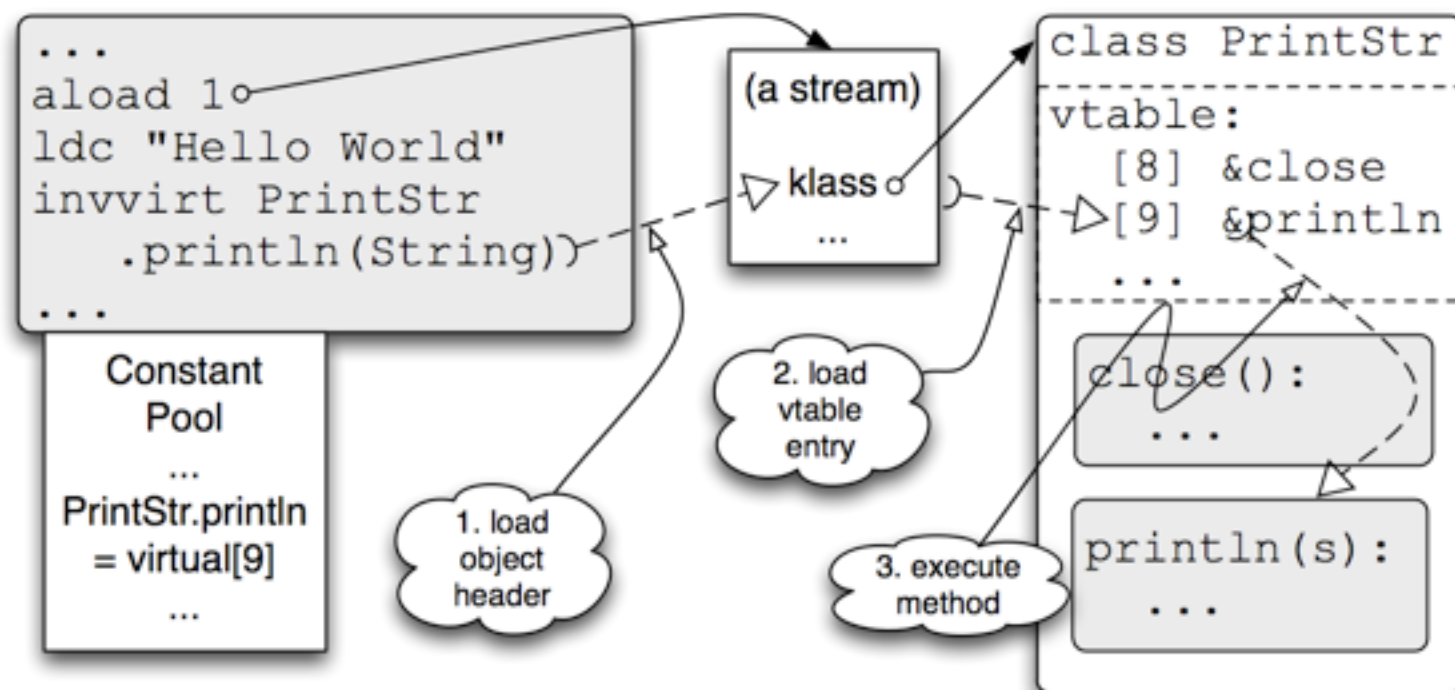
```
//PrintStream out = System.out;  
out.println("Hello World");
```

- > The compiled byte code looks like

```
4:    aload 1  
5:    ldc #2          //String "Hello World"  
7:    invokevirtual #4 //Method java/io/PrintStream.println:  
                                (Ljava/lang/String;)V
```

- a) Again, names in bytecode
- b) Again, linking fixed by JVM
- c) *Only* the receiver type determines method selection
- d) *Only* the receiver type can be adapted (narrowed)

# How the VM selects the target method:



&gt;

(Note: This implementation is typical; VMs vary.)

# What more could anybody want? (1)

- > Naming — not just Java names
  - > arbitrary strings, even structured tokens (XML??)
  - > help from the VM resolving names is **optional**
  - > caller and callee do **not** need to agree on names
- > Linking — not just Java & VM rules
  - > can link a call site to any callee the runtime wants
  - > can *re-link* a call site if something changes
- > Selecting — not just static or receiver-based
  - > selection logic can look at any/all arguments
  - > (or any other conditions relevant to the language)

## What more could anybody want? (2)

- > Adapting — no exact signature matching
  - > widen to Object, box from primitives
  - > checkcast to specific types, unbox to primitives
  - > collecting/spreading to/from varargs
  - > inserting or deleting extra control arguments
  - > language-specific coercions & transformations
- > *(...and finally, the same fast control transfer)*
- > *(...with inlining in the optimizing compiler, please)*



# Dynamic method invocation

- > How would we compile a function like

```
function max(x, y) {  
    if (x.lessThan(y)) then y else x  
}
```

# Dynamic method invocation

- > How would we compile a function like

```
function max(x, y) {  
    if (x.lessThan(y)) then y else x  
}
```

- > Specifically, how do we call `.lessThan()`?

# Dynamic method invocation (how not to)

- > How about:

```
0:    aload 1; aload 2
2:    invokevirtual #3    //Method Unknown.lessThan:
                                   (LUnknown;) Z
```

- > 5: if\_icmpeq

- > That doesn't work

- > No receiver type
- > No argument type
- > Return type might not even be boolean ('Z')

# Dynamic method invocation (slowly)

## > How about:

```
0:    aload 1; aload 2
2:    ldc #2    // "lessThan"
4:    invokestatic #3    //Method my/Runtime.invoke_2:
    (Ljava/lang/Object;Ljava/lang/Object;)Ljava/lang/Object;
5:    invokestatic #4    //Method my/Runtime.toBoolean:
    (Ljava/lang/
Object;)Z
8:    if_icmpeq
```

## > That works, but ... really ... slowly

- > Argument types are (re-)computed remotely
- > Less possibility of local caching or optimization
- > Lots of indirections!
- > (“Invoker” objects? Similar problems + complexity.)

# Dynamic method invocation (how to)

## > A new option:

```
0:    aload 1; aload 2
2:    invokedynamic #3    //NameAndType lessThan:
                                (Ljava/lang/Object;Ljava/lang/Object;) Z
>    5:    if_icmpeq
```

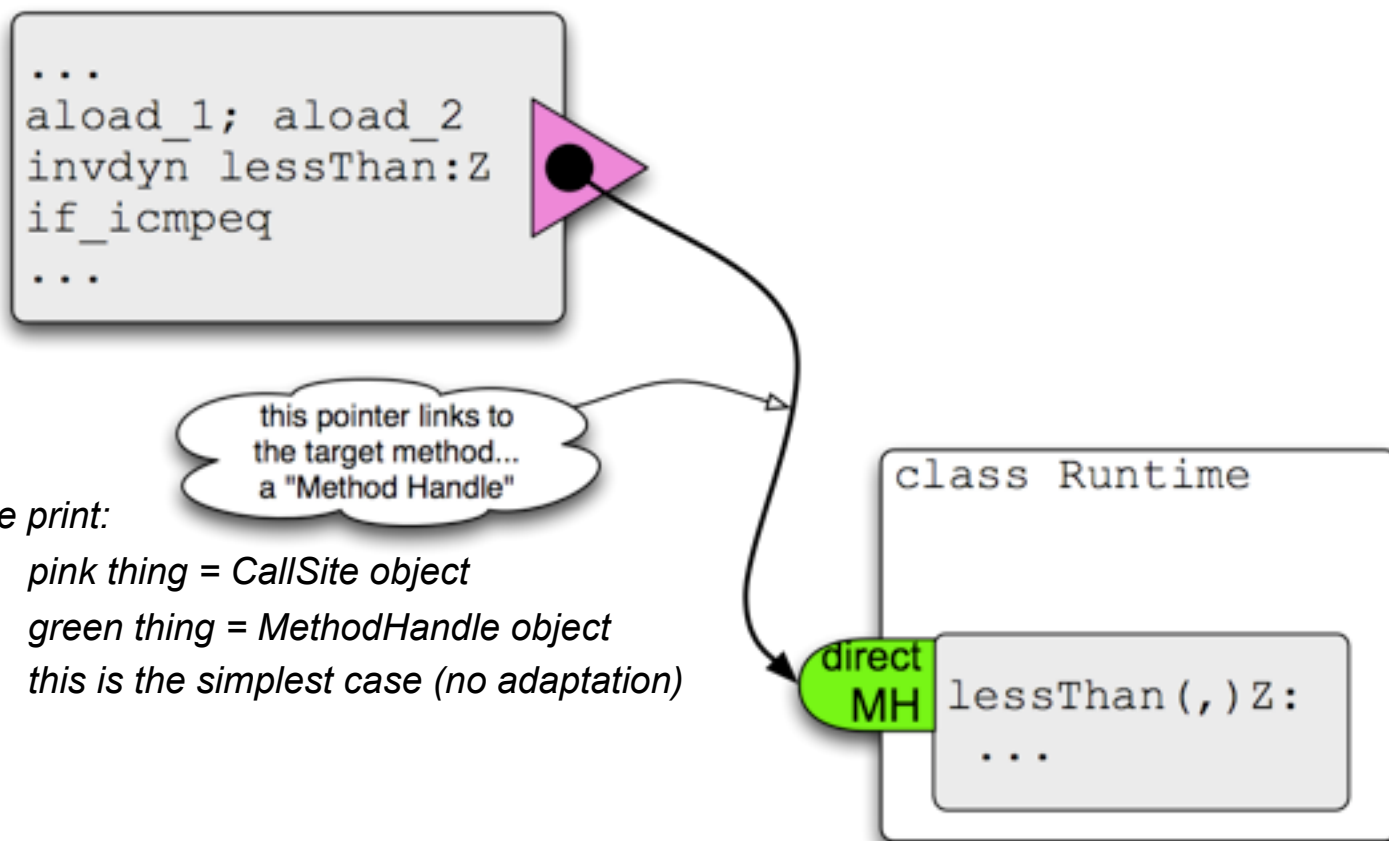
## > Advantages:

- Compact representation
- Argument types are untyped Objects
- Required boolean return type is respected
- (Flexibility from ***signature polymorphism***.)

## Dynamic method invocation (details)

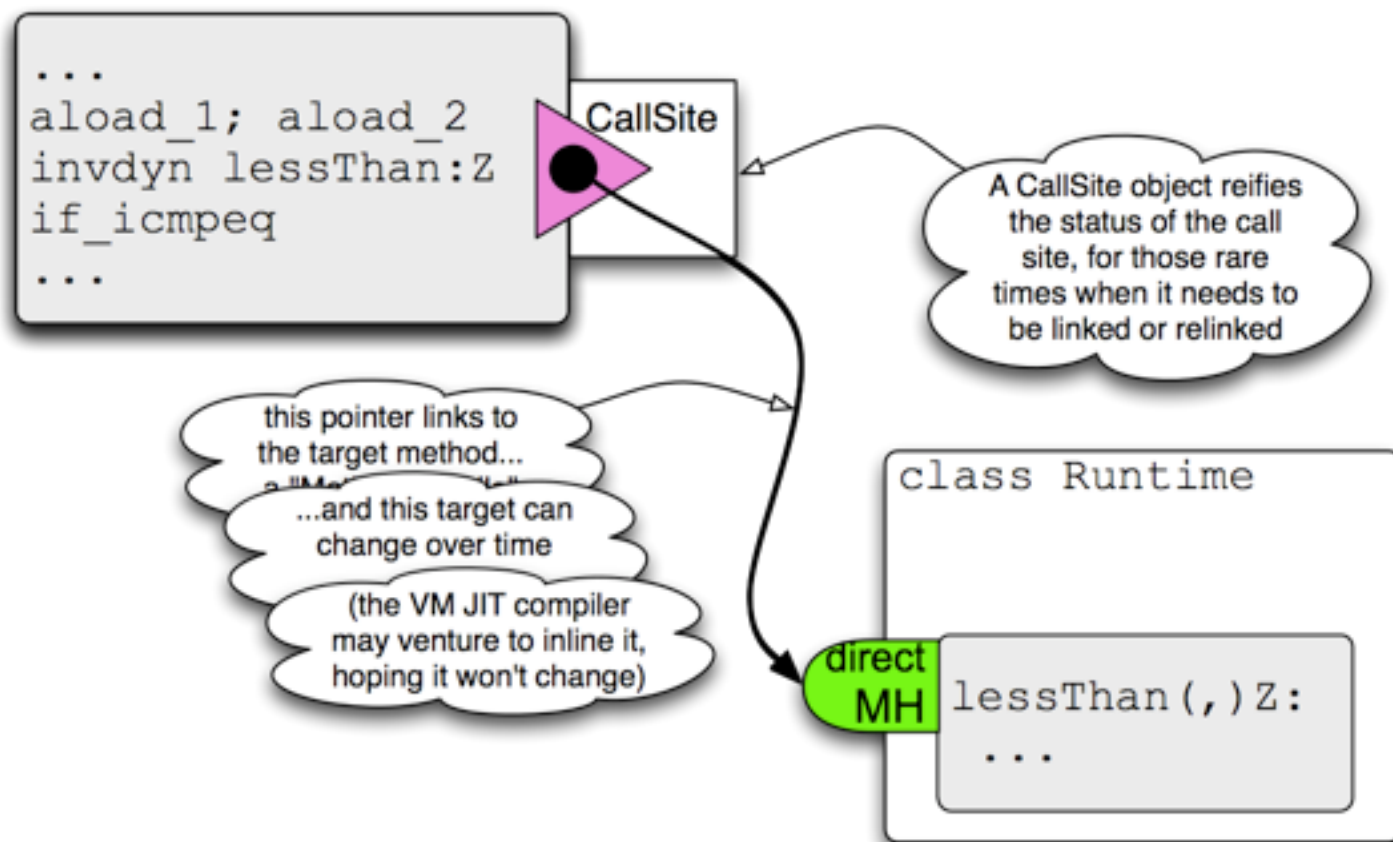
- > But where is the dynamic language plumbing??
  - > We need something like `invoke_2` and `toBoolean!`
  - > How does the runtime know the name `lessThan`?
- > Answer, part 1: it's all *method handles* (MH).
  - > A MH can point to any accessible method
  - > The target of an `invokedynamic` is a MH

# invokedynamic, as seen by the VM:



- > *Fine print:*
  - > *pink thing = CallSite object*
  - > *green thing = MethodHandle object*
  - > *this is the simplest case (no adaptation)*

# invokedynamic, as seen by the VM:





# What's this method handle thing?

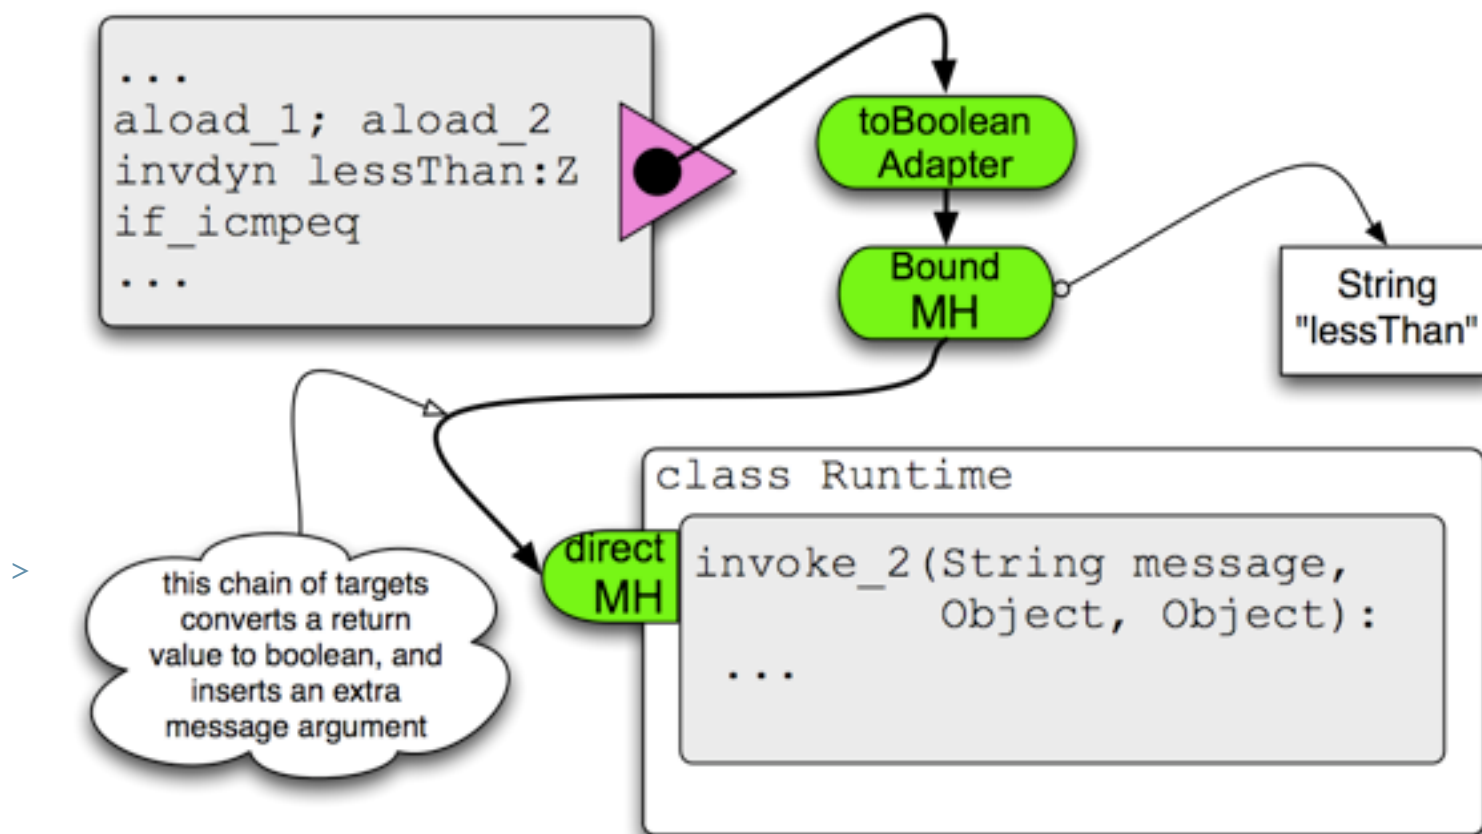
- > Fundamental unit of “pluggable” behavior
  - > Directly supports linkage to any method.
  - > (Can also do normal receiver-based dispatch)
  - > Works with *every* VM signature type.
    - > (Not just `(Object...)` => `Object` )
  - > Can be ***composed*** in chains (like pipelines)



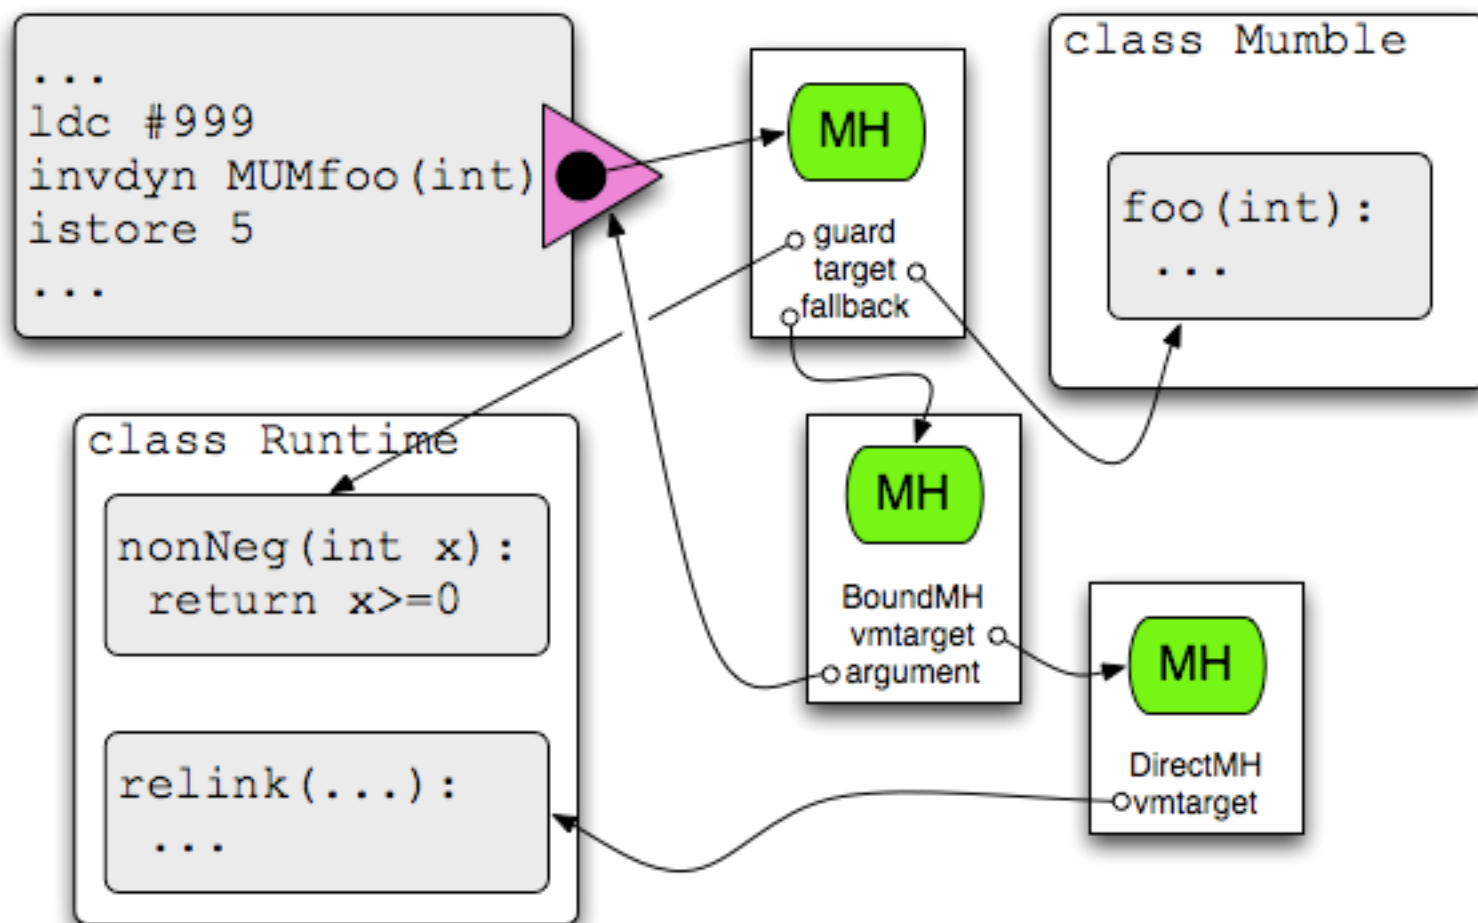
## Dynamic method invocation (details 2)

- > But can I adjust names, types, arguments?
- > Answer, part 2: In short, **yes**.
- > Method handles can be constructed, combined, and invoked in all kinds of useful patterns.
  - > Direct MH points to one method (maybe virtual)
  - > Adapter MH can adjust argument types on the fly
  - > Bound MH can insert extra arguments on the fly
  - > “Java MH” can run arbitrary code during a call
- > MHs can be chained (and the chains can be inlined)

## more invokedynamic plumbing:



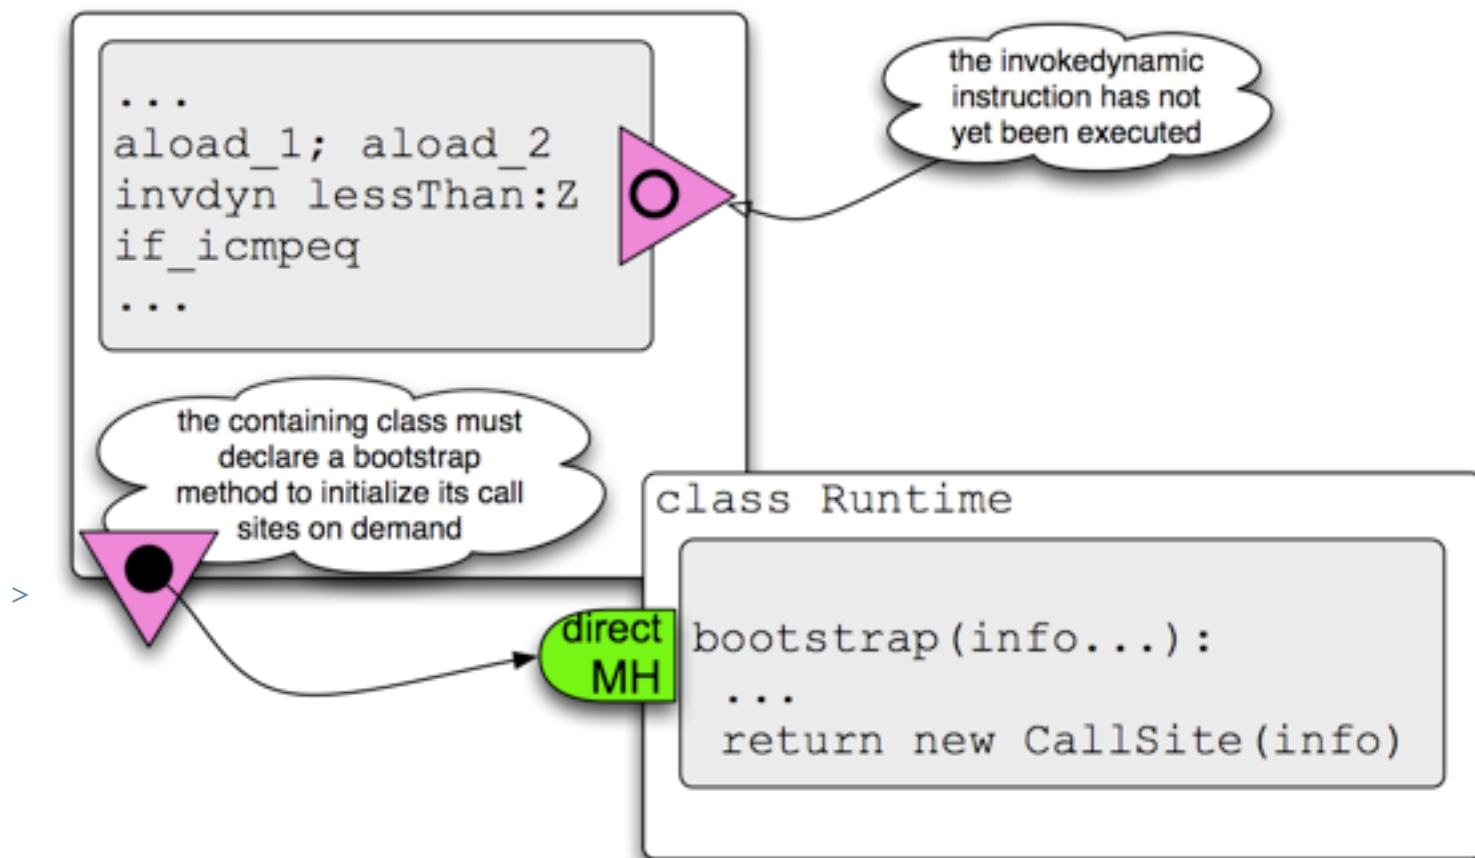
# Guarded specialization (inline cache)



## Dynamic method invocation (details 3)

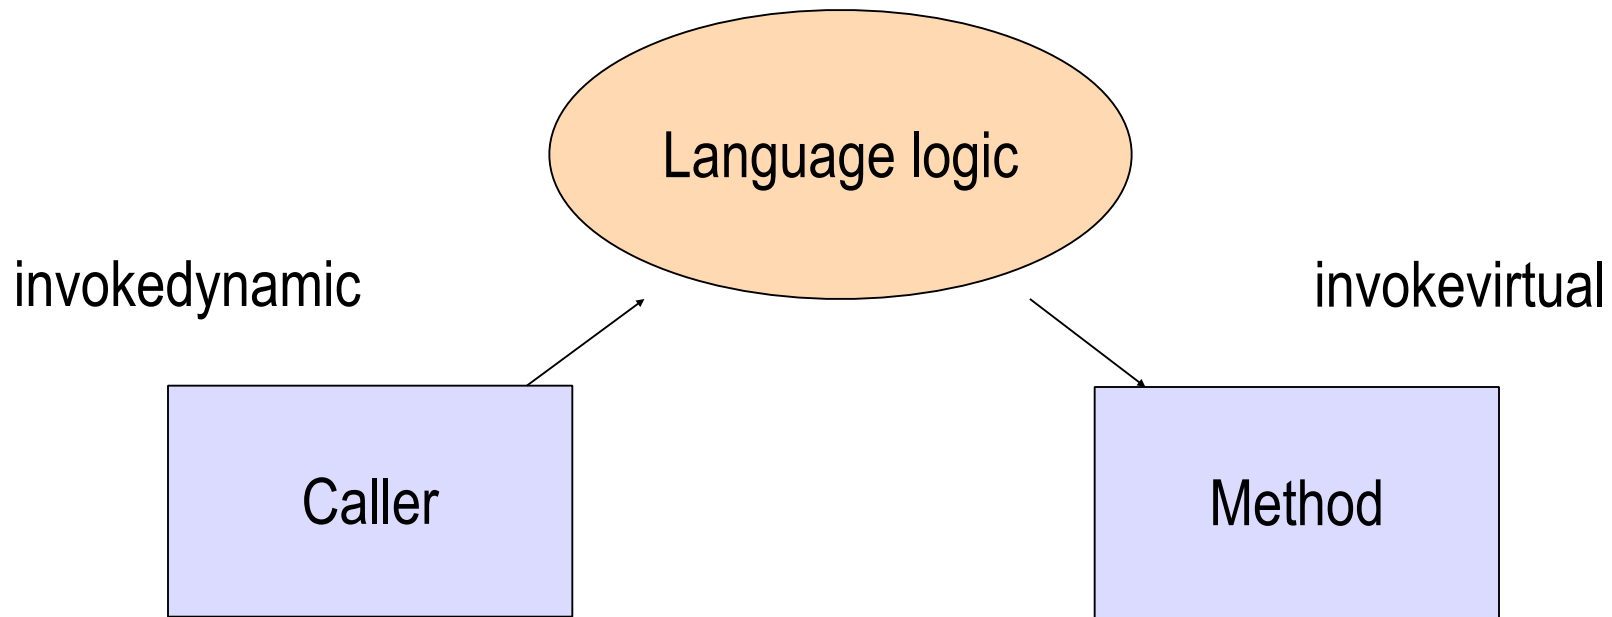
- > But who sets up these nests of method handles?
  - > We need a hook for the runtime to set them up
  - > And it needs to be a localizable, modular hook
- > Answer, part 3: a *bootstrap method* (BSM).
  - > Classes containing `invokedynamic` declare BSMs
  - > BSM is called the first time a given instruction runs
  - > The BSM gets to see all the details (name & type)
  - > The BSM is required to construct call-site plumbing
  - > (Yes, it too is a method handle.)

# invokedynamic bootstrap logic:




# A budget of invokes

<b>invokestatic</b>	<b>invokespecial</b>	<b>invokevirtual</b>	<b>invokeinterface</b>	<b>invokedynamic</b>
no receiver	receiver class	receiver class	receiver interface	no receiver
no dispatch	no dispatch	single dispatch	single dispatch	custom dispatch
<b>B8 <i>nn nn</i></b>	<b>B7 <i>nn nn</i></b>	<b>B6 <i>nn nn</i></b>	<b>B9 <i>nn nn aa 00</i></b>	<b>BA <i>nn nn 00 00</i></b>



- Check which methods are available *now* in each class [open classes]
- Check the dynamic types of arguments to the method [multimethods]
- Rearrange and inject arguments [optional and default parameters]
- Convert numbers to a different representation [fixnums]





# Optimizing invokedynamic

- Don't want it to be slow for the first 10 years
- Most of the current optimization framework applies!
- Every invokedynamic links to *one* target method so inlining is possible
- Complex language logic can also be inlined
- A JVM implementation can even assist with *speculative* optimizations that are language-specific

# Advantages of functions

- Natural unit of “live” behavior
- Simple to interpret (object method call)
- Composable (no reweaving required)
- Local expressiveness  $\approx$  well-structured bytecode
- Simple to compile (because a simple, pure value!)

# Why not always use functions?

- Not as compact as (byte-)code
- Not similar to machine code (semantic gap)
- Not clearly serializable (specific to VM instance)
- Hard to express contextual patterns (frame-relative)

## A third way: What about AST?

- Serializable but not compact
- Easy to interpret, hard to compile
- Expresses *only* contextualized patterns
- Use bytecodes for macro-ops & contextuials
- Use functions for module composition

# Resources

- John Rose (JSR 292 spec lead)
  - > <http://blogs.sun.com/jrose>
- Multi-Language Virtual Machine OpenJDK project
  - > <http://openjdk.java.net/projects/mlvm>
- JVM Language Summit, September 2009
  - > <http://www.jvmlangsummit.com>
- “JVM Languages” Google Group
  - > <http://groups.google.com/group/jvm-languages>

- JVM Language Summit,  
September  
2008, 2009

<http://www.jvmlangsummit.com>

