

Sal/Svm—A Language and Virtual Machine for Computing with Non-Enumerated Sets



Outline


- Background and Example
- Architecture
- Implementation
- Preliminary Evaluation
- Summary

Research Context and Background

- **Research context:** Hardware architectures for future device technologies
 - Are microprocessors the best way to take advantage of tradeoffs in $\lll 20\text{nm}$ CMOS?
- **Research problem**
 - **Challenge:** Can we represent compute problems independent of algorithms for their solution?
 - **Observation:** Can represent problems using set-theoretic properties their solutions obey
- **This talk**
 - Platform that represents compute problems using set-theoretic constructs and machine state

Sets and Non-Enumerative Set Representations

- Most people are familiar with these:



- Generally, people often think of sets as these:

$$S4 = \{1, 3, 7, 9\}$$

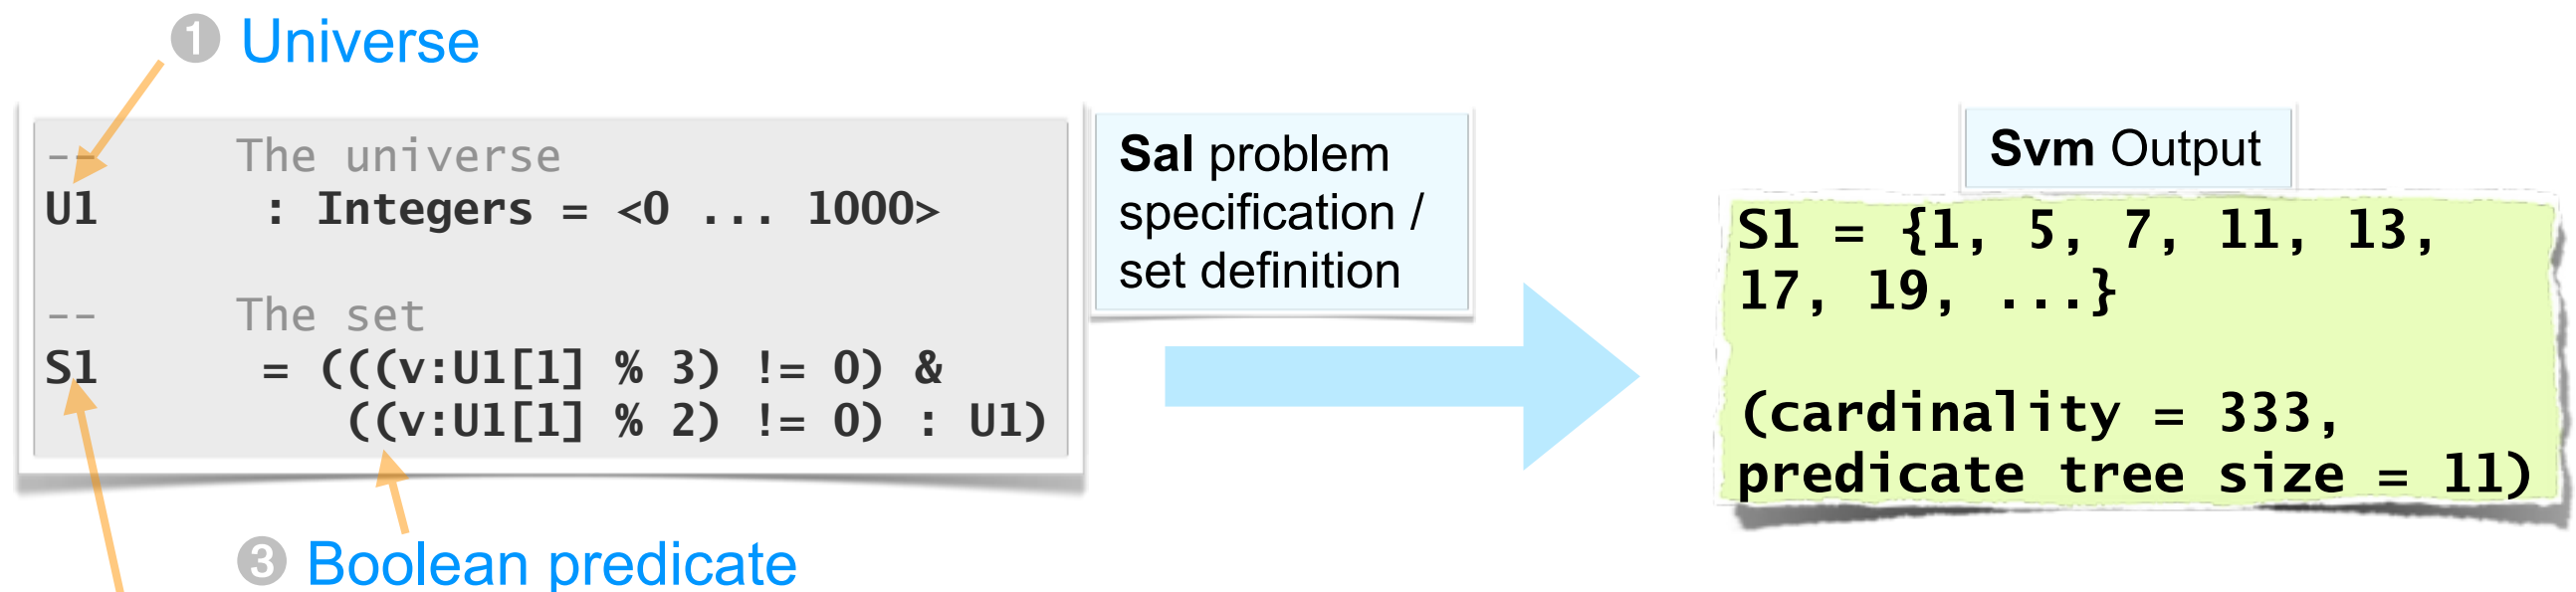
$$S5 = \{(1, \text{"science"}), (2, \text{"history"})\},$$

etc.

- Two main ways of representing sets
 - Enumerated: items in set captured in the representation, e.g., $\{1, 3, 5, \dots\}$
 - Non-enumerated: properties of the elements captured in representation: e.g. “odd integers”

Sal (Set Assembly Language) and Svm (Set VM)

- **Sal/Svm**: A Language and Virtual Machine for Computing with Non-Enumerated Sets
 - **Application**: precisely representing compute problem definitions without algorithms
 - Problems specified in an input language (Sal) for describing set relations
 - A virtual machine (Svm) processes the set relations, returning a solution (if requested / possible)
- **A simple example**: set of odd integers between 0 and 1000 that are not multiples of 3:

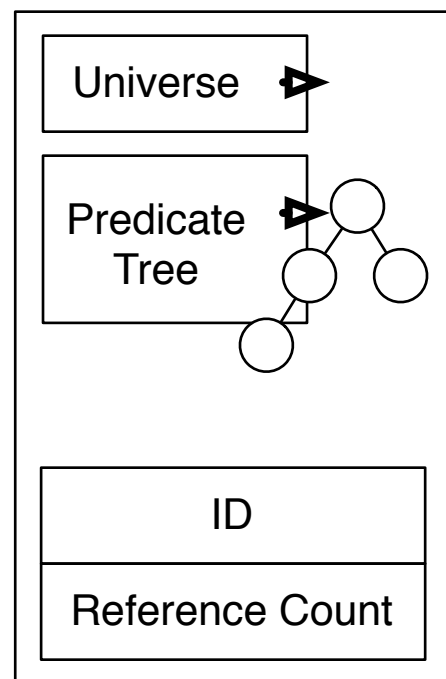


Outline

- Background and Example
- Architecture
- Implementation
- Preliminary Evaluation
- Summary

Representation of Sets in Svm

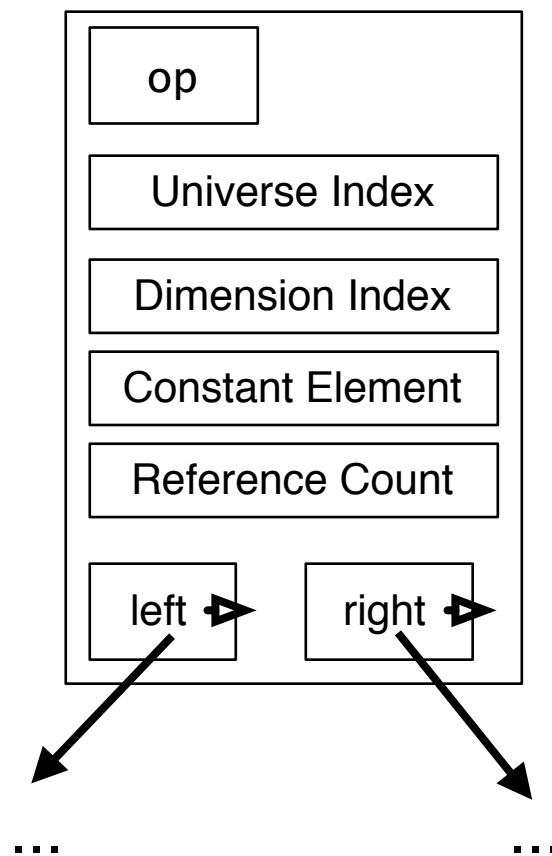
Set



=

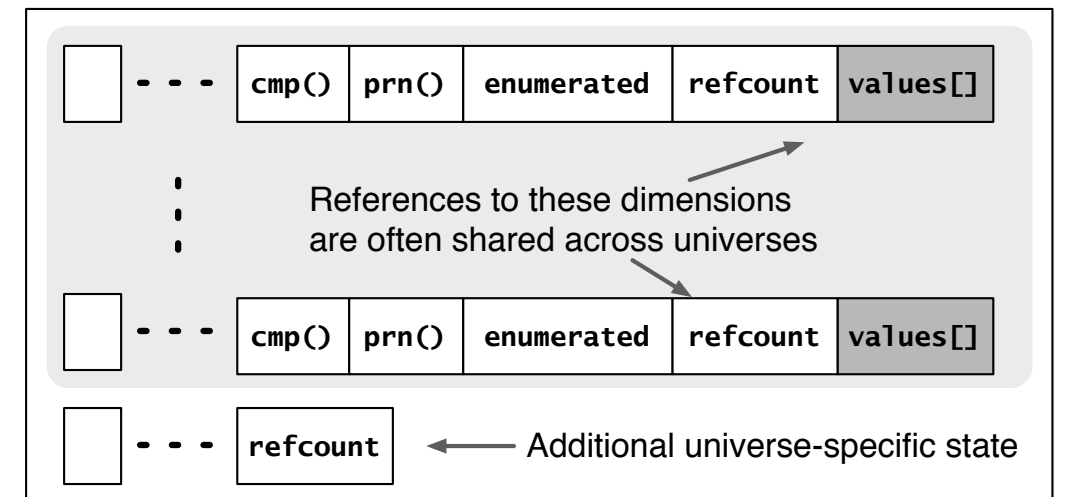
Predicate Tree

(root node shown here)



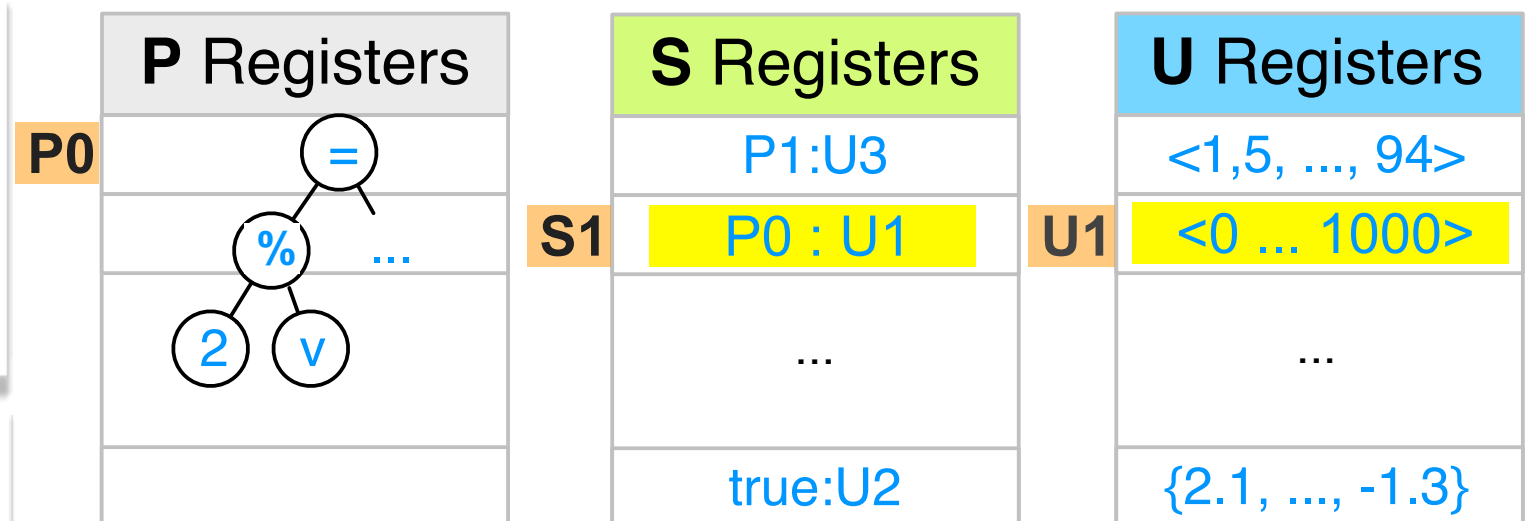
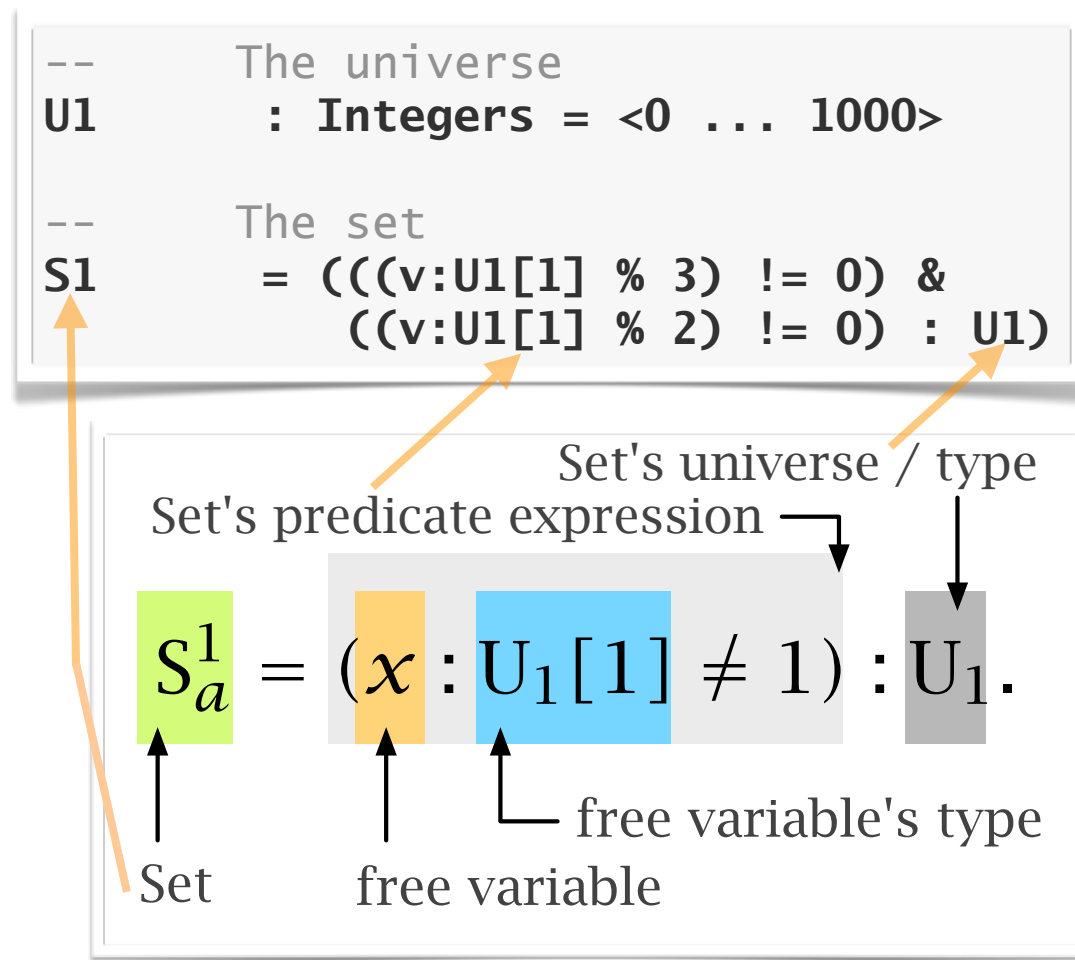
+

(Multi-Dimensional) Universe




Svm Machine Registers

- Svm virtual machine has **three sets of registers that underlie all computation**
 - **U (universe) registers:** hold the “basis” elements (scalar)
 - **P (predicate) registers:** hold Boolean predicate trees (data structures)
 - **S (set) registers:** hold information about specific pairings of predicates to universes



Larger Example: Convex Hull Compute Problem Definition



Definition (Convex Hull). The convex hull, $CH(S)$, of a set S of points in the plane, is the smallest convex polygon for which each point in S is either on the boundary thereof, or in its interior.

● Convex hull computational problem specification in Sal

```

U0      : Integers = <1 ... 10 delta 2*iota>
U1      : Integers = <1 ... 10 delta (2*iota)+1>
U2      = U0 x U1

P10     = !((qy == py:U2[2]) & (qx == px:U2[1]))
P11     = ((qx*ry - qy*rx) - px:U2[1]*(ry - qy) + py:U2[2]*(rx - qx)) >= 0
P1      = exists qx:U2[1] exists qy:U2[2] forall rx:U2[1] forall ry:U2[2] (P10 & P11)

S2      = (P1 : U2)

print   enum S2

```

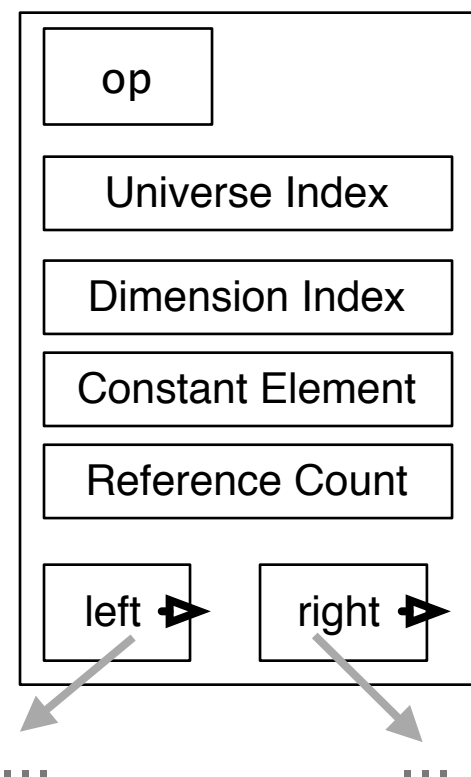
Boolean Predicate Trees

Sal predicate expression

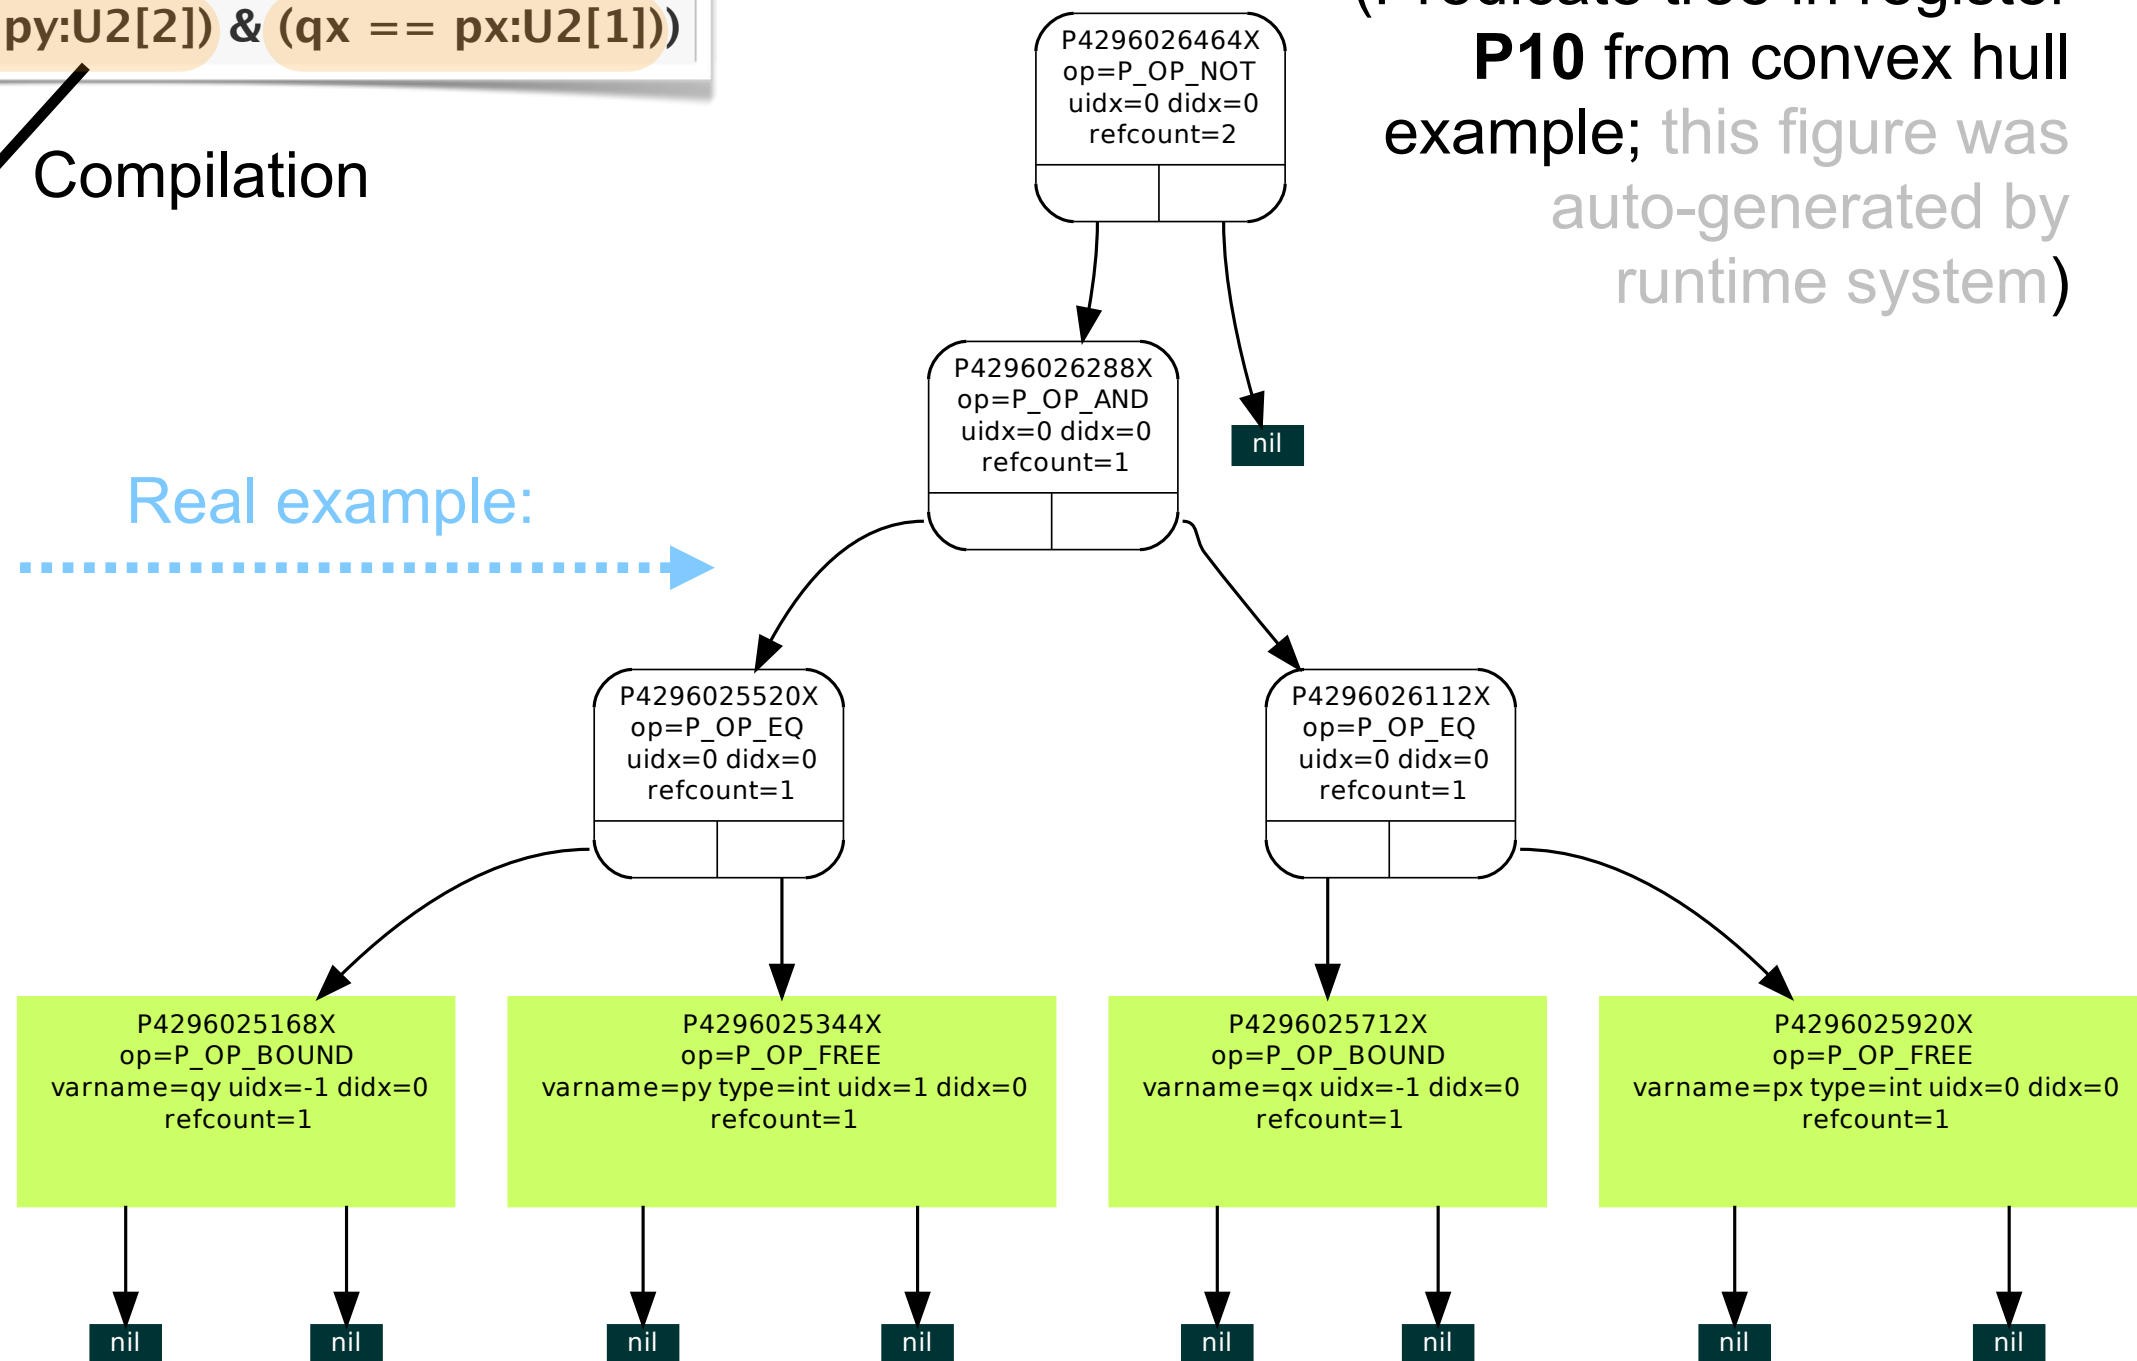
P10 = !((qy == py:U2[2]) & (qx == px:U2[1]))

Compilation

Predicate Tree



Real example:



(Predicate tree in register **P10** from convex hull example; this figure was auto-generated by runtime system)

Generated by Svm version 0.3-alpha (build 08-03-2010-20:08:47-pip@listener.local-Darwin), on Sun Aug 8 18:30:39 2010.

Outline

- Background and Example
- Architecture
- Implementation
- Preliminary Evaluation
- Summary

Implementation

● Svm

- Core of runtime system implemented in a library, in ANSI C
- Library also provides utilities such as rendering machine state (example in previous slide), etc.

● Salc


- Implemented in C, with YACC-driven parser front-end
- Takes Sal source assembler and outputs intermediate representation in an ELF container
- A few optimizations implemented, potential for more

● Interactive command console

- Accepts Sal statements, parses and injects into machine state
- Interactive help system, commands to probe Svm state, etc.

● Web interface to interactive console on server

- Runs interactive command console via web, on server



- Interactive interface is embedded in a content-management system
 - Contains examples that can be modified and run via web; users can post comments, questions

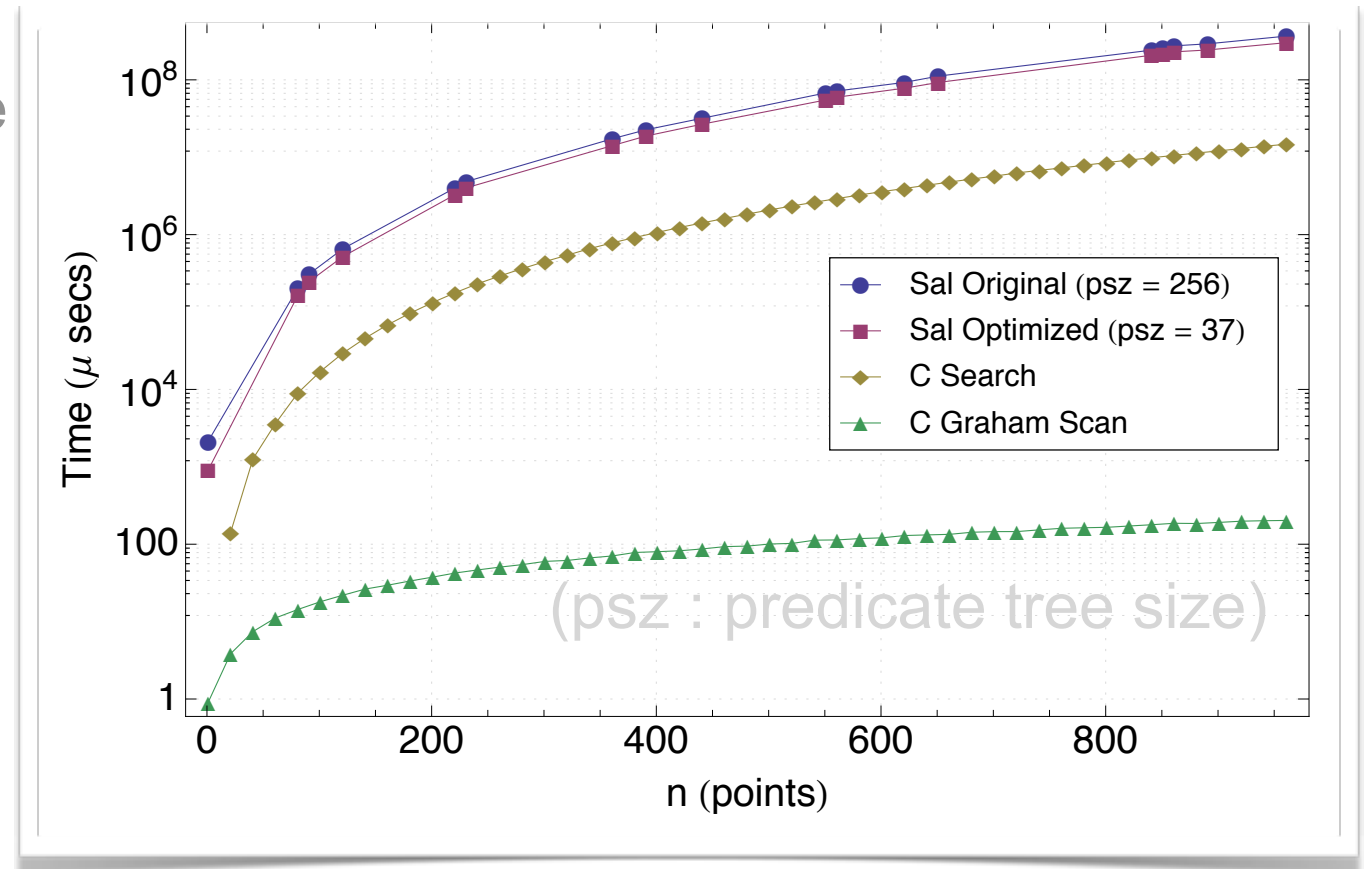



Outline

- Background and Example
- Architecture
- Implementation
- Preliminary Evaluation
- Summary

Preliminary Performance Evaluation: Convex Hull Problem

Recall: given set of points in plane, compute “points on the hull” (black points below)



● Evaluation setup

- All cases: compiler: **gcc 4.2.1**, -O3; CPU: **2.8GHz Intel® Core™ i7**; OS: **MacOS® 10.6.3**

● Svm is within ~20× of optimized C code implementing same enumeration algorithm

- Sal **problem specification** is significantly simpler (~11 lines) than C **algorithm implementation**

● Orders of magnitude slower than (provably) optimum algorithm (Graham's Scan)

- **Paper provides a counterexample where Sal/Svm and C search can be more energy-efficient**

Outline

- Background and Example
- Architecture
- Implementation
- Preliminary Evaluation
- Summary

Summary and Ongoing Work

- Summary: **Research context**
 - Architectures for future devices and device technologies
- Summary: **Sal/Svm**
 - A **framework for specifying computational problems** using set-theoretic constructs
 - Enables leaving algorithm choice undefined, to pick best hardware-dependent tradeoff
- Ongoing: **Integrating Sal-based problem definitions in high-level languages**
 - Ongoing work to use Sal problem definition as part of type structure in a high-level language
 - Specification of “arithmetic imprecision” in computation problems
- **Availability**
 - Currently available to users inside IBM Research
 - Working to make interactive web version available to general public in coming months

