



# **ADVENTURES IN MICRO- BENCHMARKING**

**Yossi Gil, Keren Lenz, Yuval Shimron**  
**The Technion**

# WHY IS MICRO-BENCHMARKING INTERESTING

Because, deep inside, we still believe that the whole is the sum of its parts:

What would be the impact of replacing a component **A** by component **B** in a program **P**?

Compare  $P_A$  with  $P_B$ ?

Why not compare **A** and **B**? The results would be then valid for all **P**!

# WHY IS MICRO-BENCHMARKING DIFFICULT?

- The CPU / memory model / operating system /JIT / virtual machine / ... stack is very complicated
  - No one fully understands it all.
- Microbenchmarking can never be done directly.
  - A java method may consume less than 10 ns of CPU time!
- Exponential blowup in the parameters space:
  - CPU model / # cores / Memory Size
  - Software environment: O/S, etc.
  - JVM
  - JVM parameters
- JVM peculiarities:
  - Multi-threading (13 threads in a “thread free” program)
  - Garbage collection
  - JIT optimization

# BACKGROUND I

## ANONYMOUS ECOOP PAPER

- **Berlin 2007:** Timing results (hidden in tables and in slides):
  - $X$ : time to process a data structure of size  $n$
  - $Y$ : time to process a data structure of size  $10n$
  - But  $X < Y/10$
- **Berlin 2007 (lunch break):** *“I can show you the results right here, check it out on my laptop...”*
- **Berlin 2007 (after lunch break):** *“I will email you the results...”*
- **Portland 2011:** ...

# BACKGROUND II

## SPACE OPTIMIZATION OF JAVA'S HASH MAP

- **Anonymous Reviewer:** *“I do not believe your timing results... Did you account for garbage collection cycles? Did you allow sufficient time for the JIT to become active? Did you experiment with different VM flags?”*
- **Author response:** *“I am confident that the results are correct!”*
- **Rejection Aftermath:** *“There must be a system in this madness!”*
- **Well known authority:** *“Billions and billions of runs!”*

## BACKGROUND III: THE LAW OF LARGE NUMBERS

$$\Pr\left(\lim_{n \rightarrow \infty} \frac{1}{n} \sum_{i=1}^n x_i = \mu\right) = 1$$

- Let  $x_1, \dots, x_n$  be a sequence of measurements of some phenomenon (a random variable). Then, with probability 1, the average of the sequence converges to the expected value of the phenomenon.
- Sounds trivial... But, very powerful.

No matter what, if you repeat your measurement sufficiently many times, and then average over all measurements, you will get the “right” result.

## BACKGROUND IV: THE CENTRAL LIMIT THEOREM

$$\lim_{n \rightarrow \infty} \Pr \left[ \sqrt{n} \left( \frac{1}{n} \sum_{i=1}^n x_i - \mu \right) \leq z \right] = \Phi \left( \frac{z}{\sigma} \right)$$

- Let  $x_1, \dots, x_n$  be a sequence of measurements of some phenomenon .
- Then, the average of the sequence, the deviation from the expected value follows a **normal distribution**.
- Further, the standard deviation of this distribution is inversely proportional to the square root of the sequence length.

# SO, WHY SHOULD ANYTHING GO WRONG?

## REPEAT

- Experiment:
  - A. Define a simple microbenchmark
  - B. Neutralize all “noise” factors: GC, JIT, background processors, Bill Gates, etc.
  - C. Carefully read Josh Bloch’s warnings
  - D. Use billions and billions of runs.
- Plot/Tabulate/Blah Blah the results

UNTIL “happiness achieved”;



# OUR WAY OUT OF THIS INFINITE LOOP

## REPORT DIFFICULTIES ENCOUNTERED

- **Instability of the Virtual Machine**

- Different invocations give statistically different results
- Disabling the JIT makes things worse

- **The steady state rules:**

- I. The “steady state” will be different in different invocations
- II. The “steady state” may suddenly change during the same invocation
- III. At any given time during an invocation may have more than one steady state

- **Long Memory of the Virtual Machine:**

- Even short execution may contaminate the JIT
- So, what’s the point in microbenchmarking?

# EXPERIMENTAL SETTING

- **Hardware \*** : Intel Core 2 Quad CPUQ9400, 8GB
- **Java**: OpenJDK, IcedTea6 1.9.8,JVM 1.6.0\_20
- **O/S**: Ubuntu \* 10.04.2 \*
- **Run mode**: single user (`telinit 1`), text mode (no GUI), clean boot, no network, batch execution,wait for small uptime before starting
- **Benchmarked code \*** : function `get` in the JRE's standard collection class `HashMap`
  - Bit operations: rotate/XOR
  - Memory dereferencing
  - Comparisons
  - Conditionals/Iterations/function call / return
  - No dynamic dispatch
  - No memory allocation/dispatch

# BENCHAMRKing PROCEDURE

- Minimal \* pre-processing in the main Java program.
- Monitor each measurement using MX beans found in class `ManagementFactory`
- Discard **measurement** if:
  - GC cycle detected.
  - JIT cycle detected
  - Load/Unload event detected
- Function `getCaller()` calls `get()`  $n$  times,  $n = 1,152$
- Call function `getCaller()`  $m$  times,  $m \sim 6,000$  (total runtime is 100ms \*).
- **Measure** the time of each call
  - Note this runtime `getCaller()` is expected to follow a normal distribution.
- Repeat  $r$  times,  $r = 20,000$

# RAW RESULTS OF SEVEN INVOCATIONS

No	Mean	SD	Median	MAD	MIN	MAX
1	15.27	0.0802	15.27	0.0509	14.95	15.69
2	15.17	0.0768	15.17	0.0454	14.74	15.69
3	15.27	0.0726	15.24	0.0451	14.93	15.67
4	15.38	0.0727	15.38	0.0430	15.06	15.81
5	15.29	0.0867	15.30	0.0495	14.95	15.68
6	15.10	0.0722	15.11	0.0420	14.83	15.51
7	15.24	0.0929	15.25	0.0559	14.87	15.60

Looks decent, right?

# NO! SOMETHING IS OBVIOUSLY WRONG!!!

Largest Average

No	Mean	SD	Median	MAD	MIN	MAX
1	15.27	0.0802	15.27	0.0509	14.95	15.69
2	15.17	0.0768	15.17	0.0454	14.74	15.69
3	15.27	0.0726	15.24	0.0451	14.93	15.67
4	15.38	0.0727	15.38	0.0430	15.06	15.81
5	15.29	0.0867	15.30	0.0495	14.95	15.68
6	15.10	0.0722	15.11	0.0420	14.83	15.51
7	15.24	0.0929	15.25	0.0559	14.87	15.60

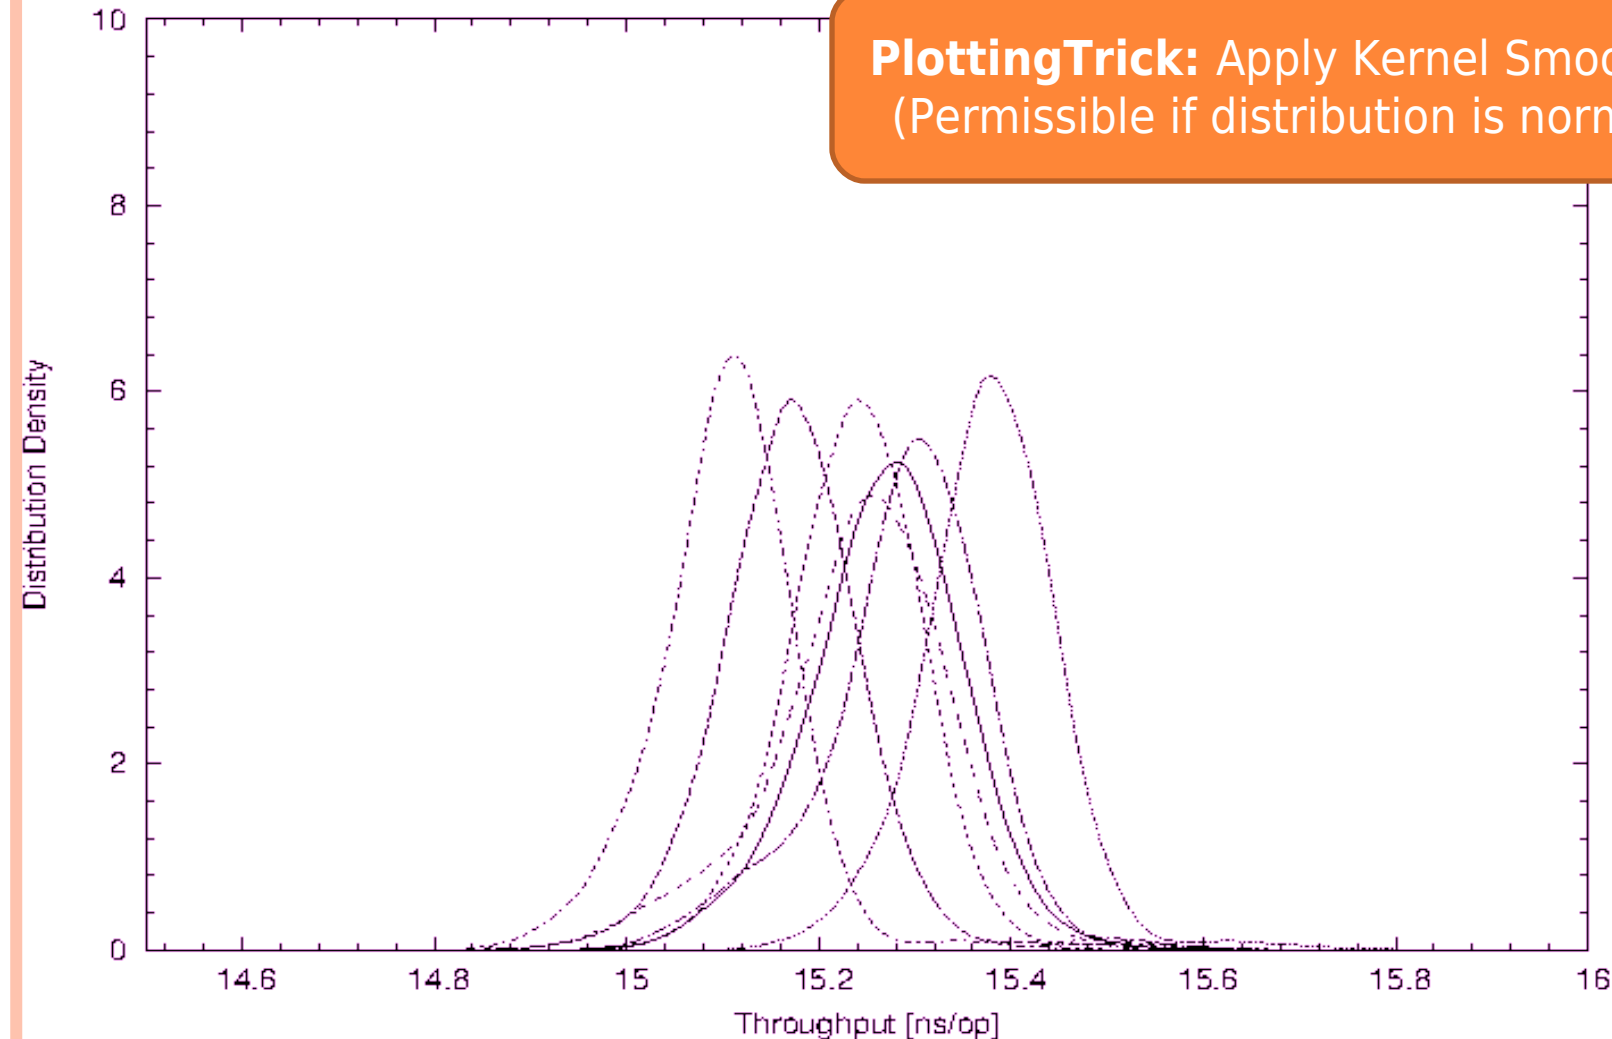
Smallest Average

Largest SD

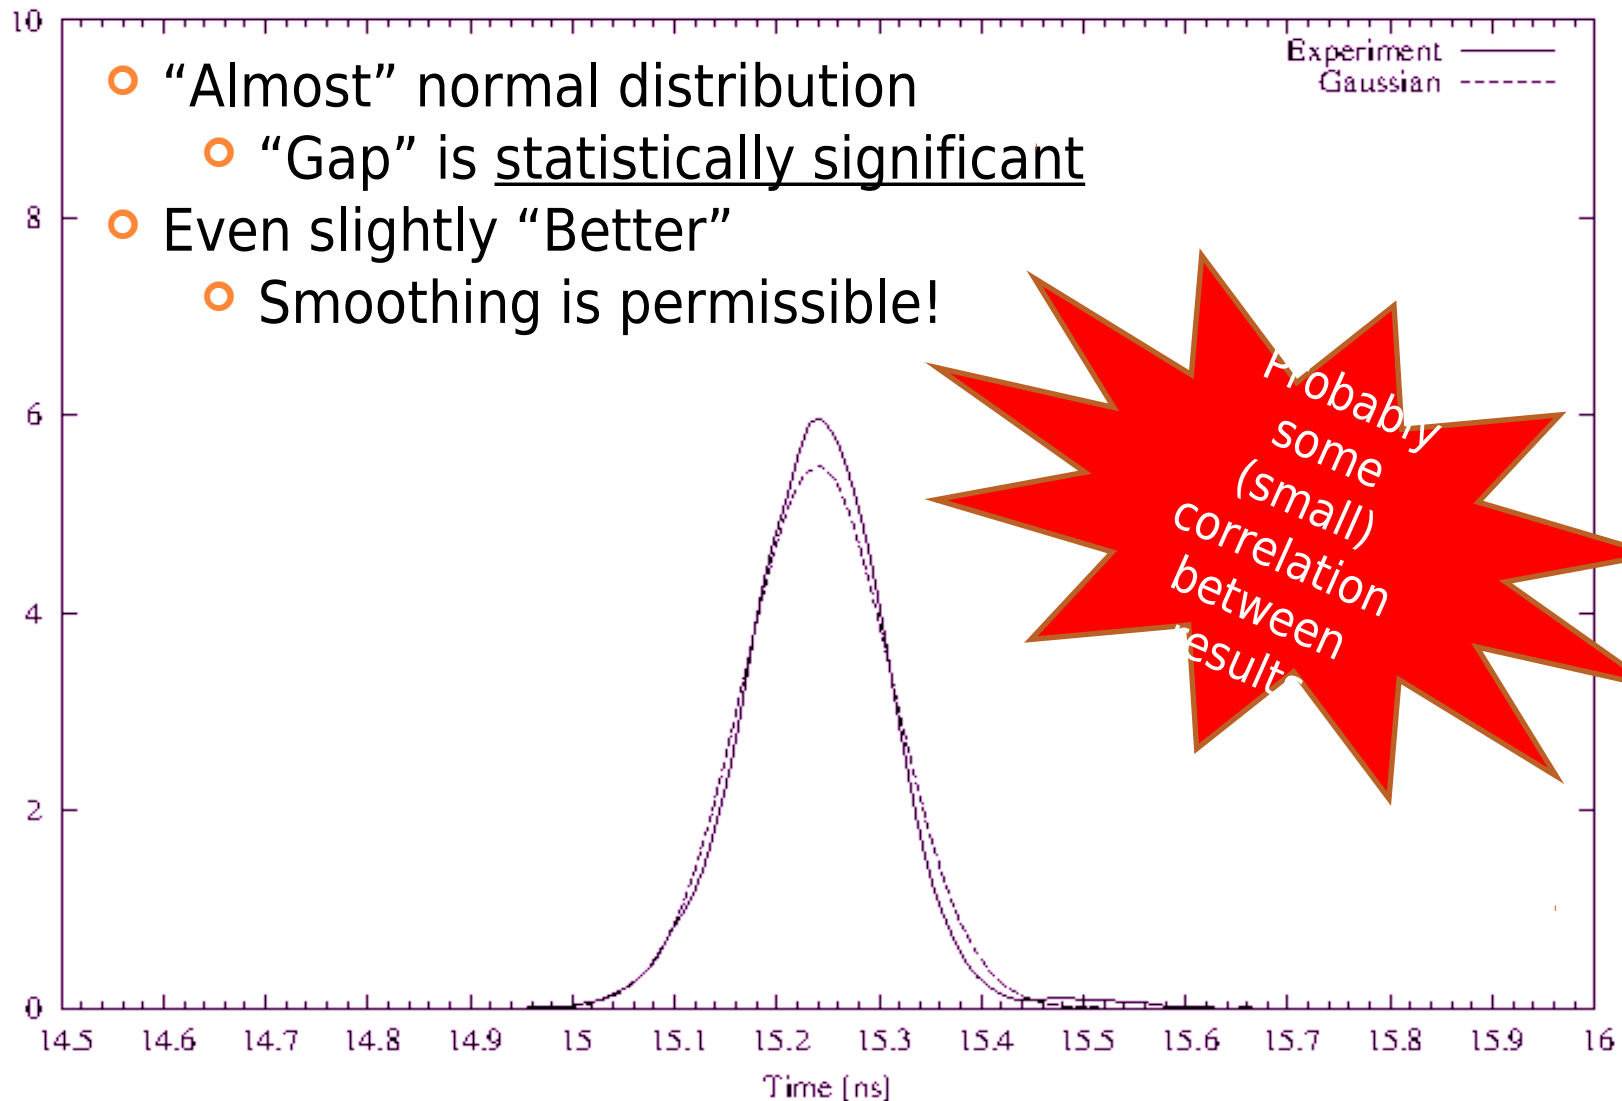
Difference is 3  
Standard  
Deviations!!!

# DISTINCT INVOCATIONS DO NOT SHARE A SINGLE STEADY STATE!


**PlottingTrick:** Apply Kernel Smoothing  
(Permissible if distribution is normal)



# ARE THE MEASUREMENT RESULTS IN A SINGLE INVOCATION NORMALLY DISTRIBUTED?




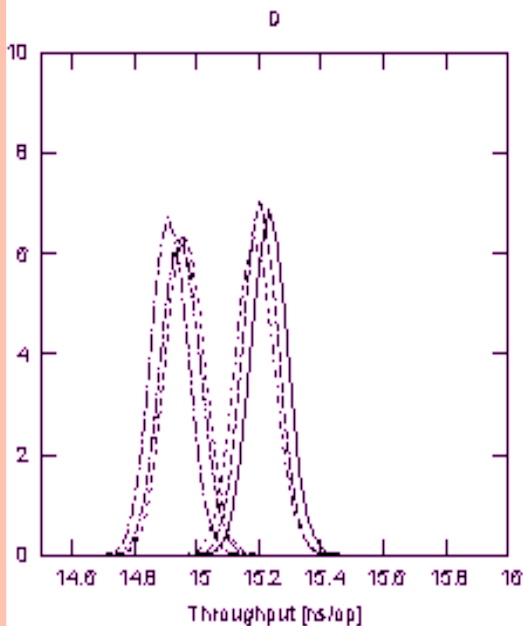

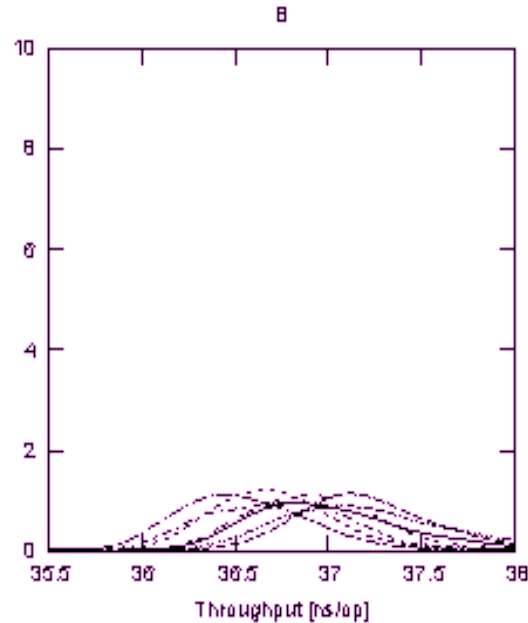
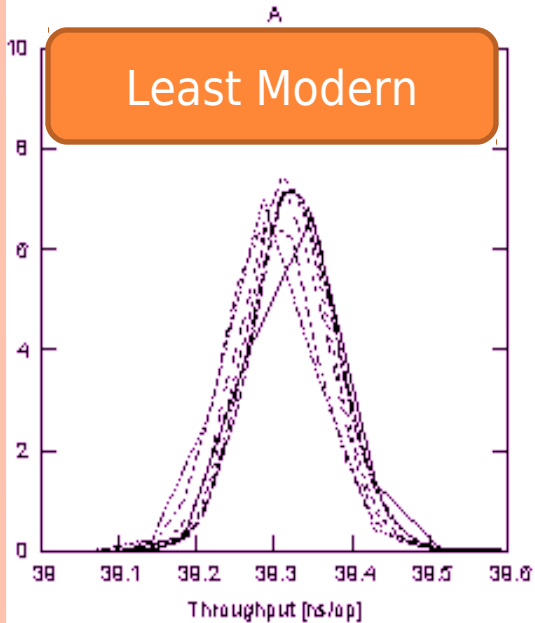
# SITUATION IS WORSE IF JIT IS DISABLED



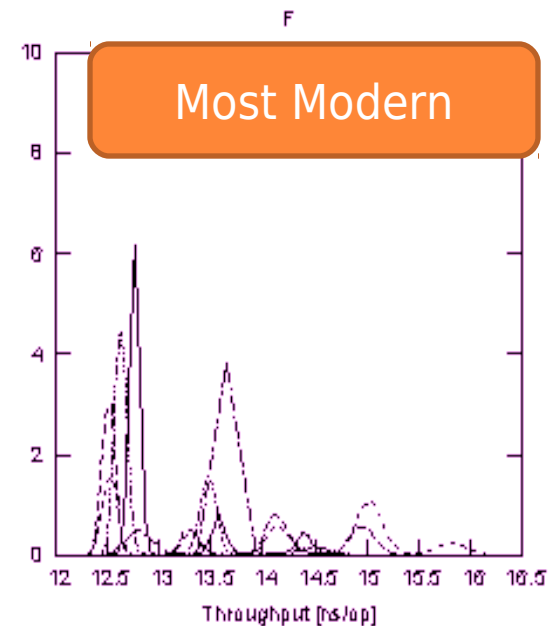


# EXPLORE OTHER COMPUTING PLATFORMS

Least Modern



Most Modern



# VARIATION ON BENCHMARKED CODE


- Variations:

- Array Bubble Sort
- List Bubble Sort
- XOR of Random Values
- Ergodic List


- Results:

- Essentially the same
- XOR-Random is slightly different:
  - Distinct peaks
  - Measurements in a single invocation are far from a normal distribution

# MULTIPLE STEADY STATE IN LIST BUBBLE SORT

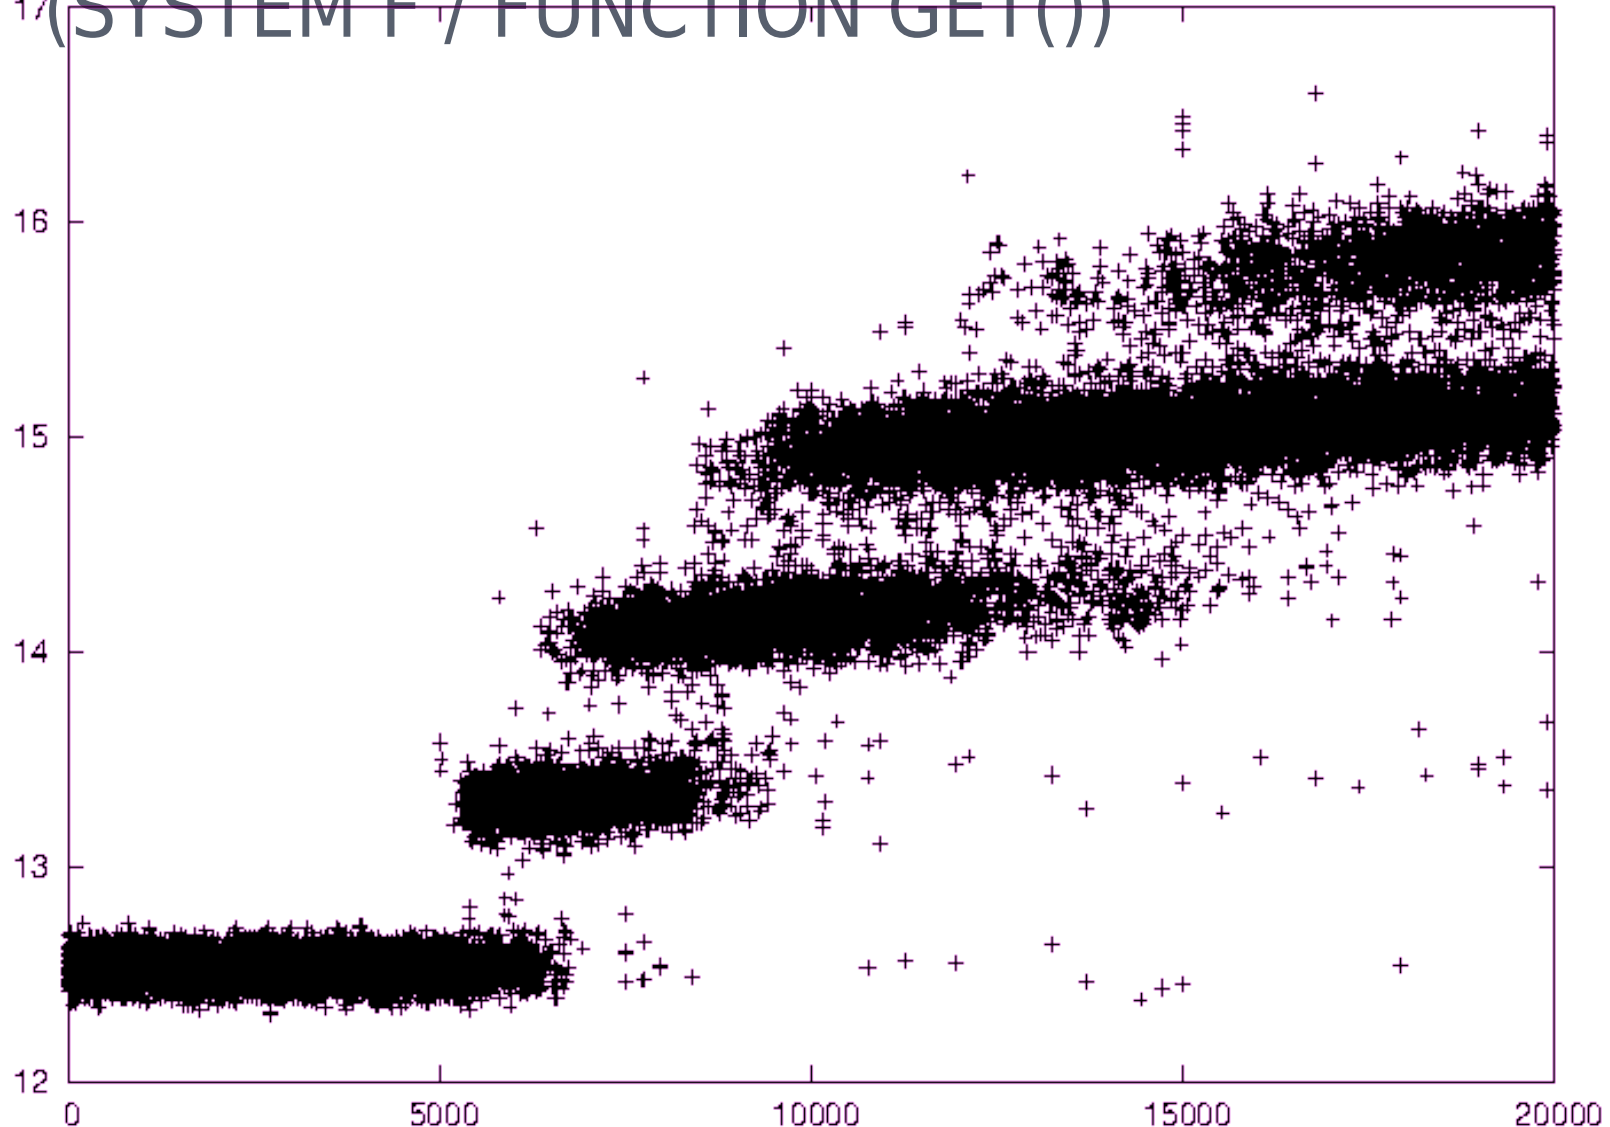


# MULTIPLE STEADY STATES IN XORRANDOMOMS




# CONCURRENT MULTIPLE STEADY STATES

(SYSTEM F / FUNCTION GET())



# NO RECOVERY FROM PRLOGUE RUNS

FUNCTION GET NEVER REACHES THE 15NS PERFORMANCE, AFTER JUST 50 FOREIGN CALLS



# CONCLUSIONS?

- Questions!
  - Did we miss some ghost threads?
  - Why?
  - Statistical methods for compensations?
  - Is the trouble of micro-benchmarking worth the dubious results?
- On the more positive side:
  - Statics is still valid: if you do many invocations, you obtain some estimate on the mean and on the SD.
  - Compilation planning is probably at fault of the JIT memory.
- Further research:
  - Spend many more months exploring the exponential parameter space.